

ID2203 - Tutorial 1

Distributed Systems, Advanced Course



Cosmin Arad

Tallat Shafaat

Seif Haridi

icarad@kth.se

tallat@kth.se

haridi@kth.se

KTH - The Royal Institute of Technology

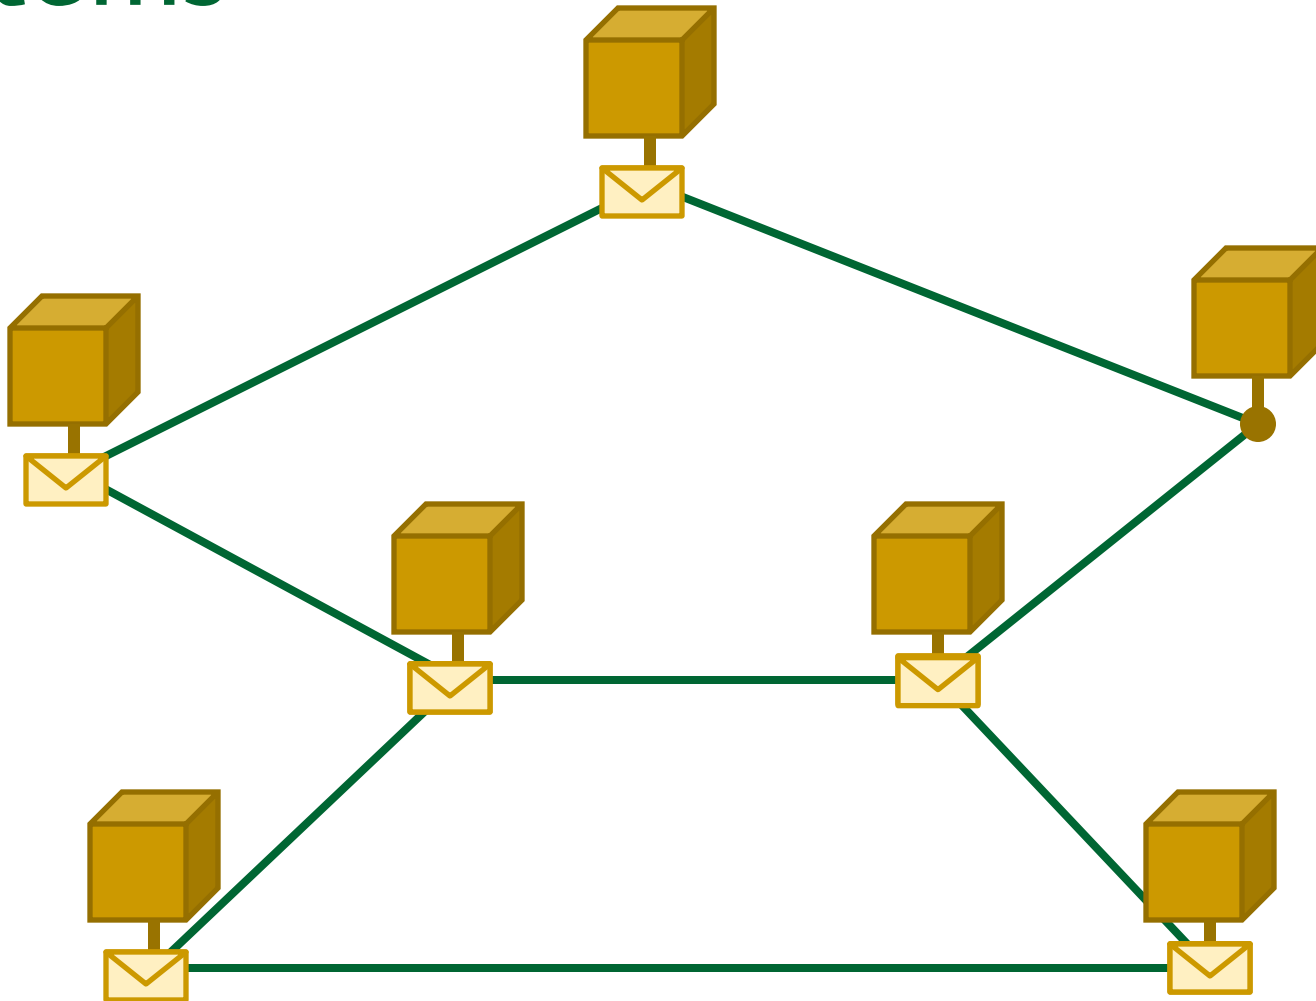
Overview

- Introduction to Kompics
- Relation to the textbook
- Assignments framework
- First assignment

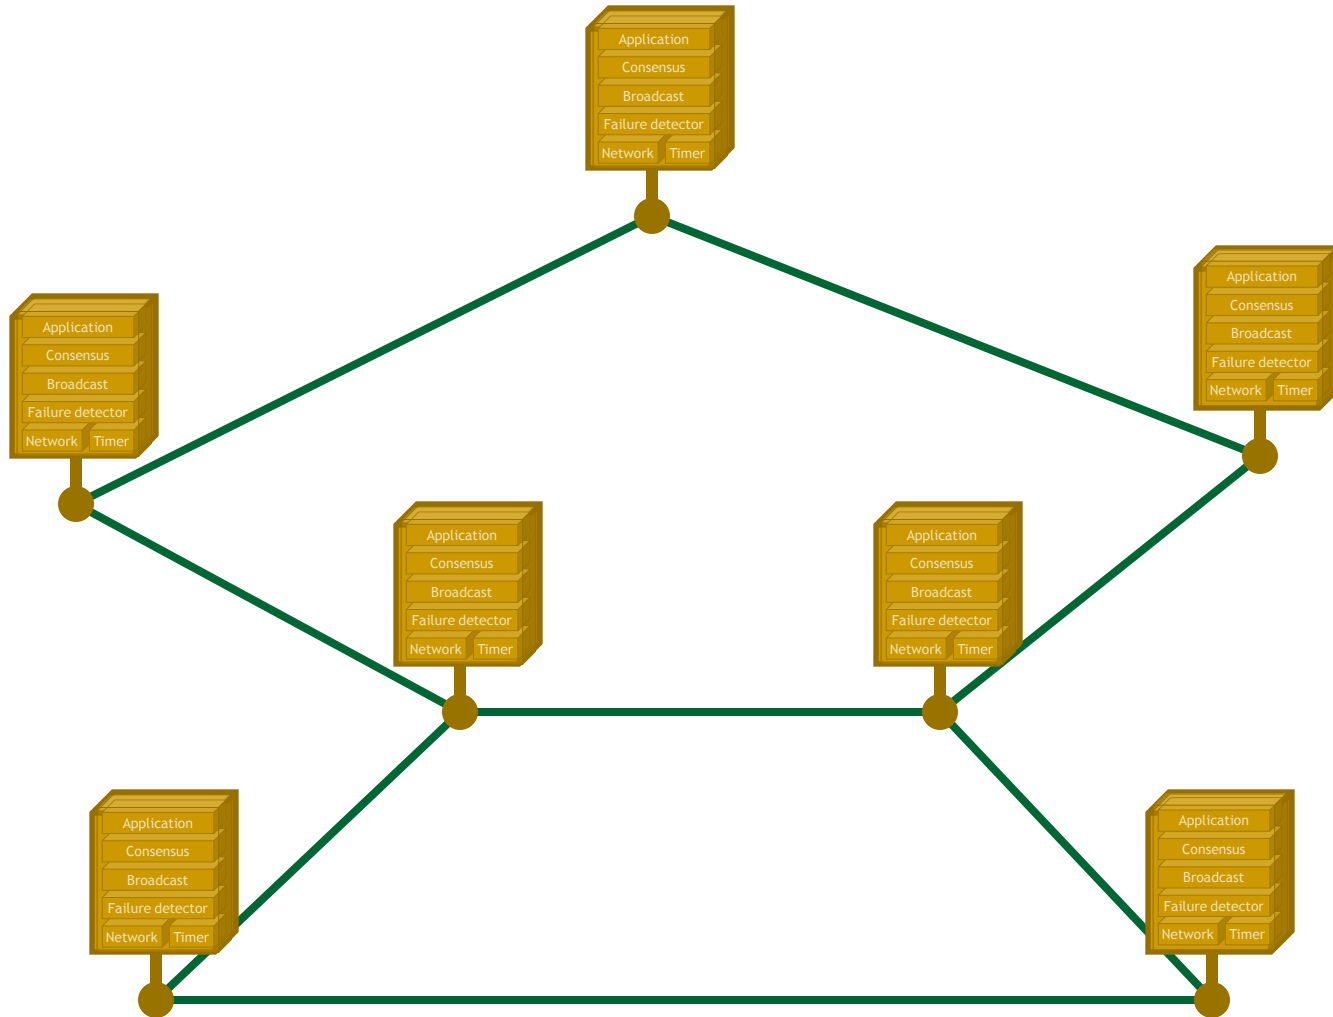


Component Model for Distributed Systems

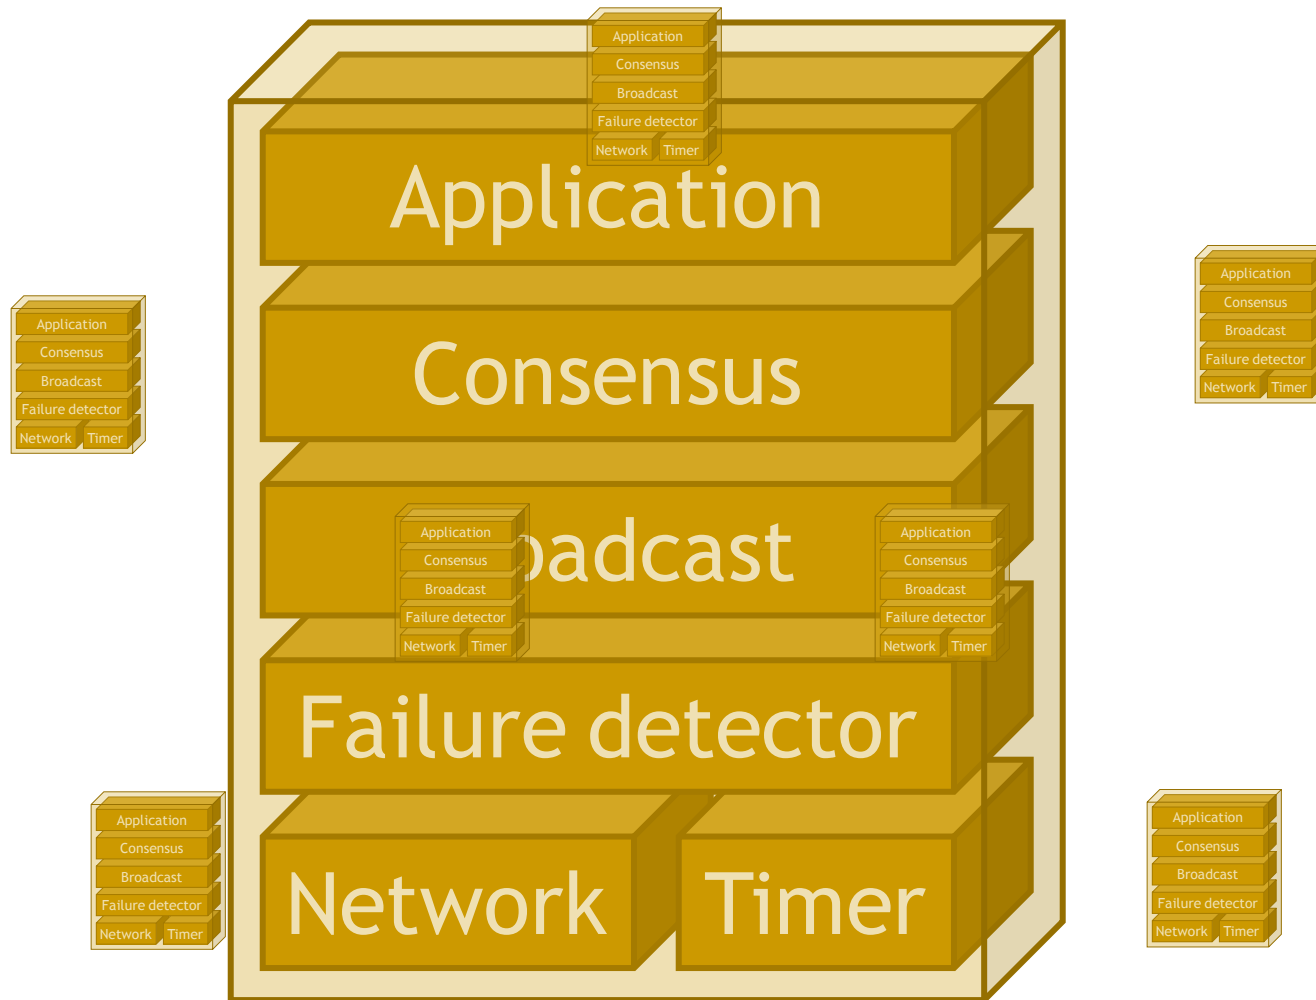
We want to build distributed systems



by composing distributed protocols



Implemented as reactive components



with message-passing concurrency



Concepts in Kompics

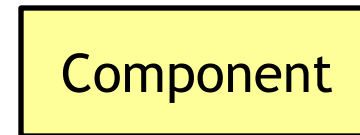
- Event



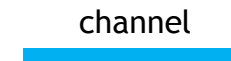
- Port



- Component



- Channel



- Handler



- Subscription



- Publication / Trigger



Events



Event

- Events are passive immutable objects
 - with typed attributes / fields
- Events are typed and can be sub-typed

```
class Message extends Event {  
    Address source;  
    Address destination;  
}  
  
class DataMessage extends Message {  
    Data data;  
    int sequenceNumber;  
}
```



Message



DataMessage

$\text{DataMessage} \subseteq \text{Message}$

Ports

Port



Direction



Direction

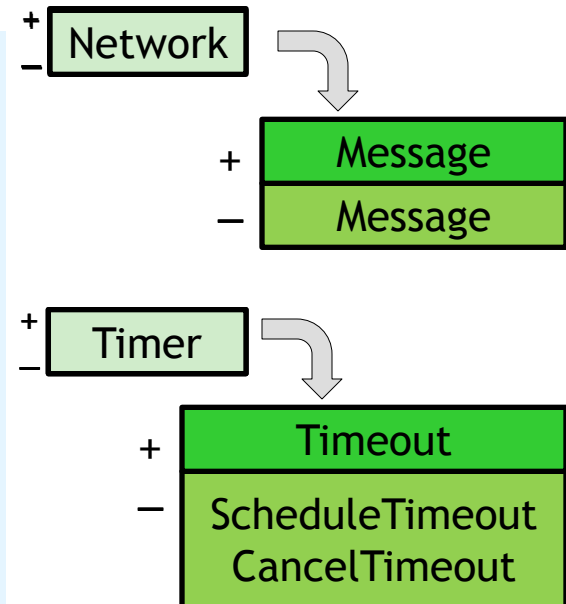
Ports

Port

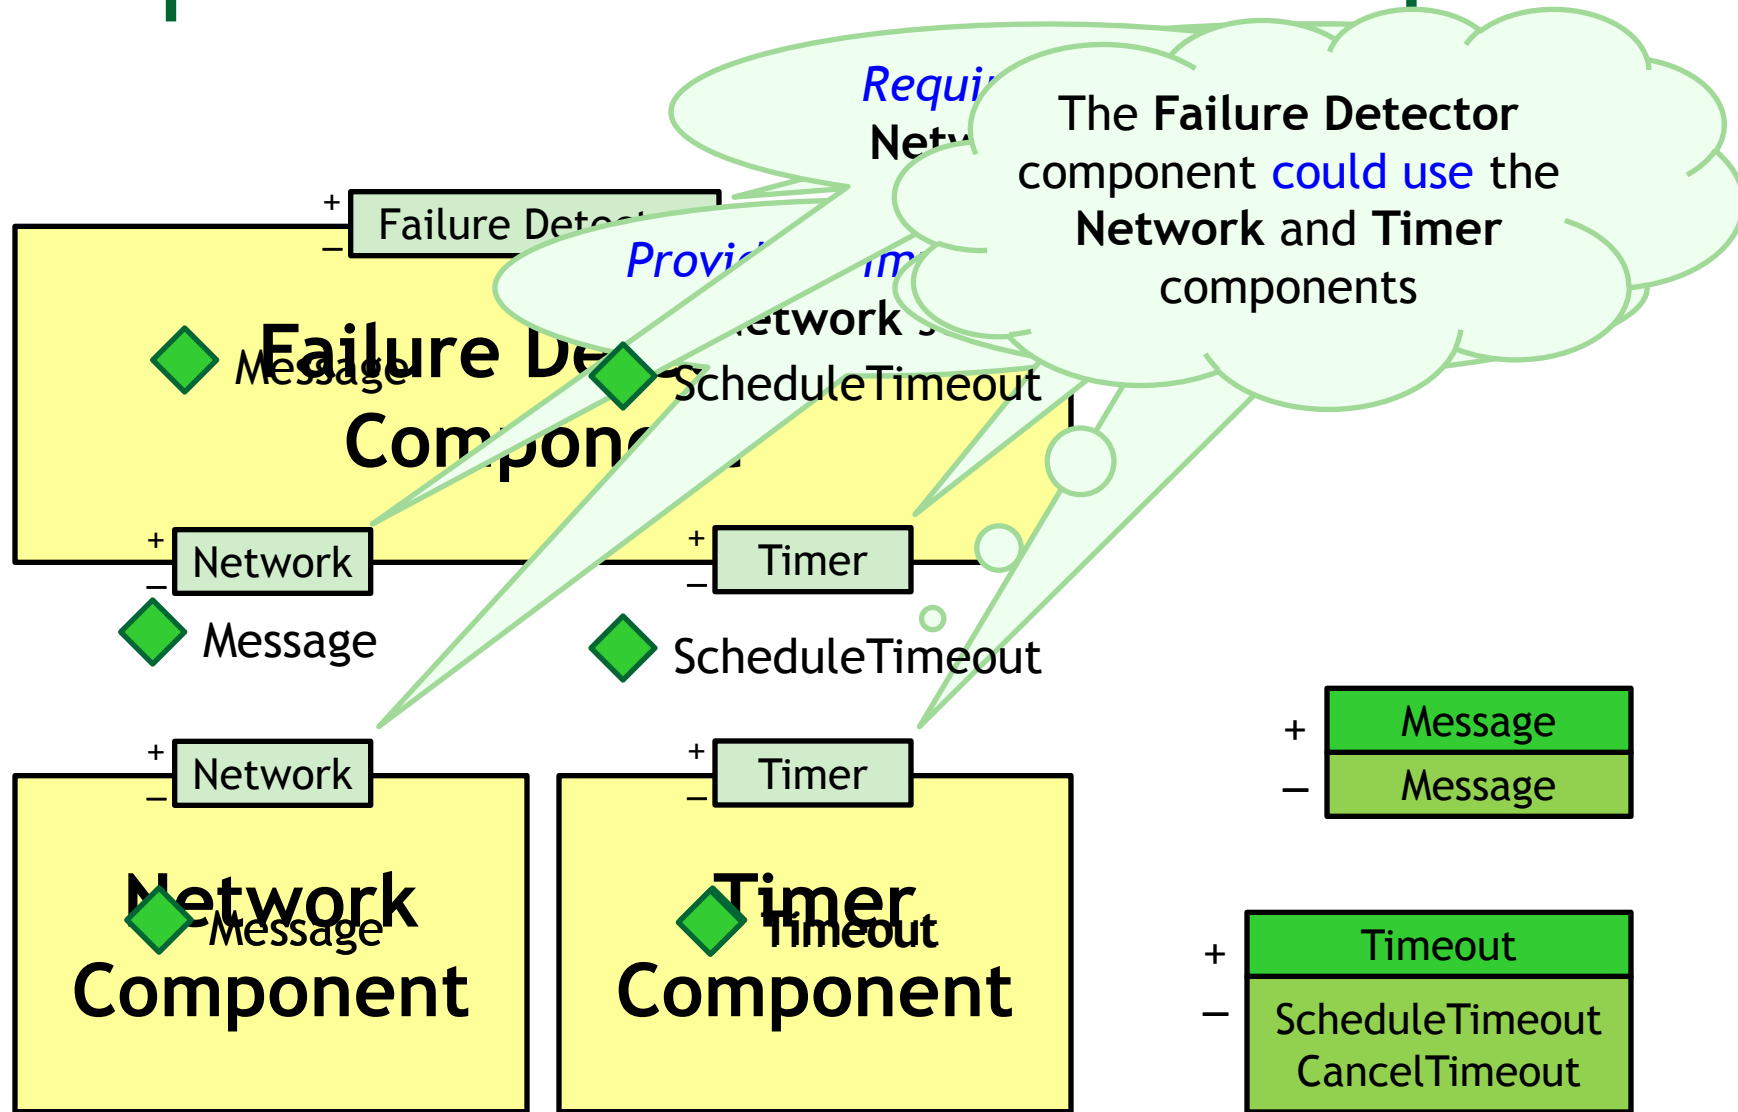
- are **bidirectional** event-based comp interfaces
 - have a **positive (+)** and a **negative (-)** direction
- A *port type* consists of 2 sets of event types
 - one set of event types for each direction, **+** and **-**
 - represents a service/protocol abstraction

```
class Network extends PortType {{
    positive(Message.class);
    negative(Message.class);
}}

class Timer extends PortType {{
    positive(Timeout.class);
    negative(ScheduleTimeout.class);
    negative(CancelTimeout.class);
}}
```



Components with Ports Example



Channels

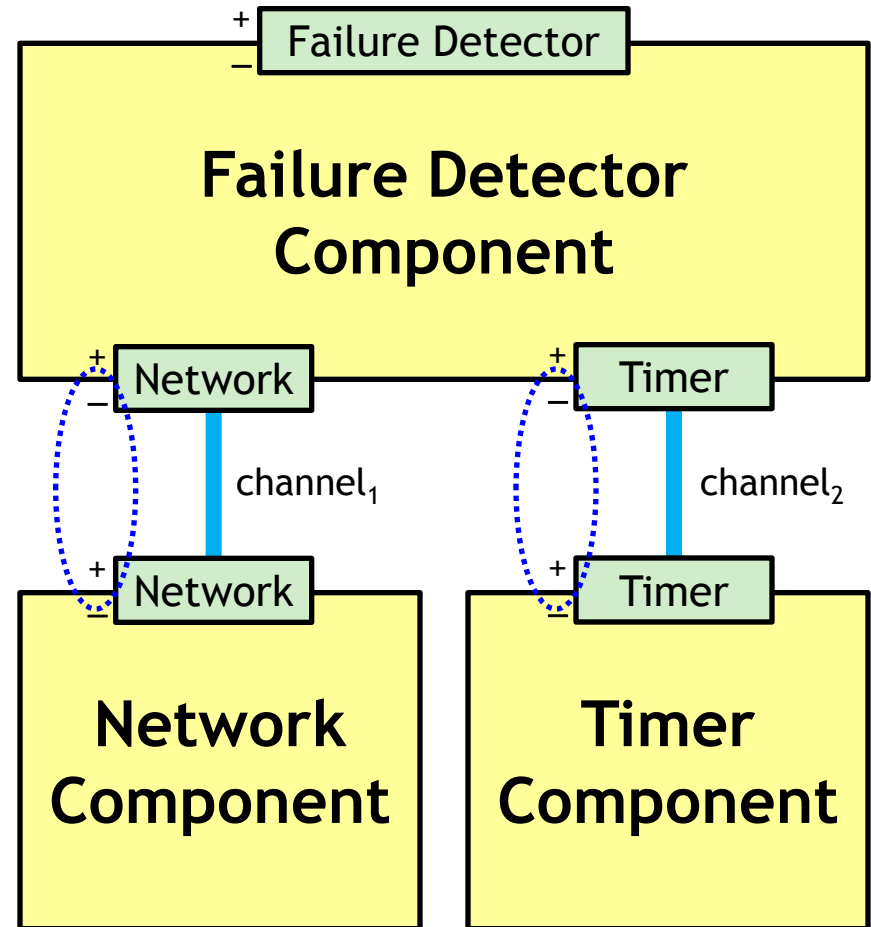
channel

- Channels connect *complementary* ports of the *same* type

□ + to -

□ - to +

- Channels forward events in *FIFO* order in *both* directions



Event handlers

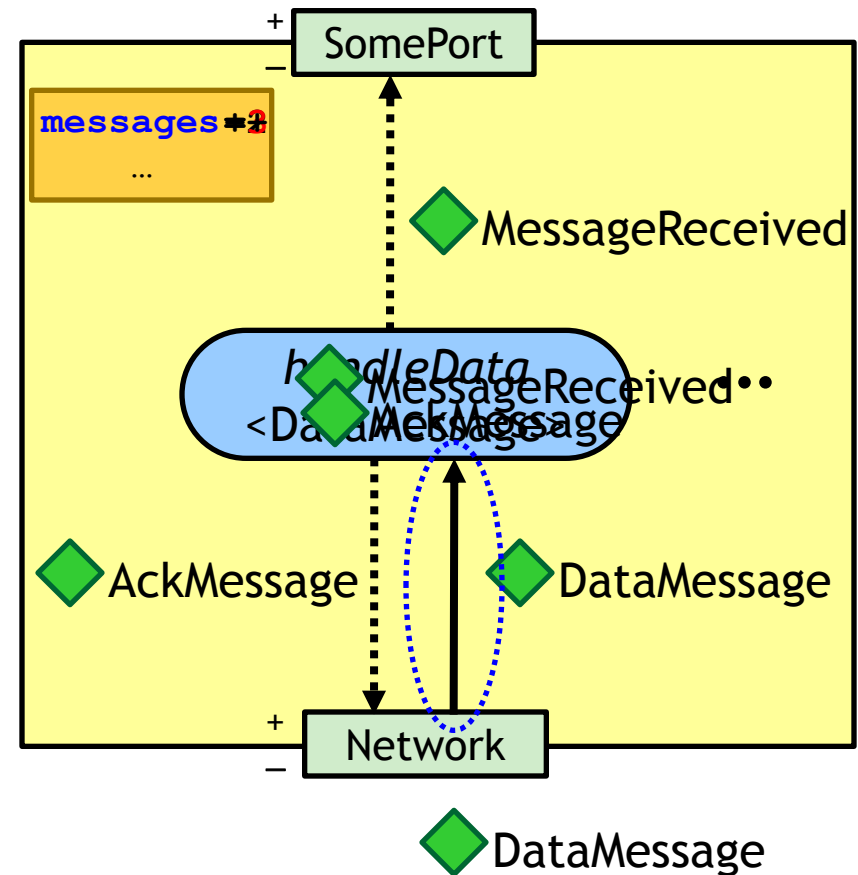
handler

- A handler is a **first-class** component procedure
- accepting a particular type of events
- executed reactively upon receiving an event
 - may mutate the local state and trigger new events
 - handlers of **one** component are **mutually exclusive**

```
Handler<DataMessage> handleData = new Handler<DataMessage>() {  
    public void handle(DataMessage dm) {  
        messages++; // component state update  
        trigger(new MessageReceived(dm.data), somePort);  
        trigger(new AckMessage(myAddress, dm.source,  
                                dm.sequenceNumber), network); // event triggering  
    }  
};
```

Subscriptions & Publications

- A subscription **binds** an event handler, h , to a local component port p
- Let e be the type of h
 - let f be a supertype of e
 - f must come in on p
- After subscription
 - h will handle all events of any type d , subtype of e , coming in on p

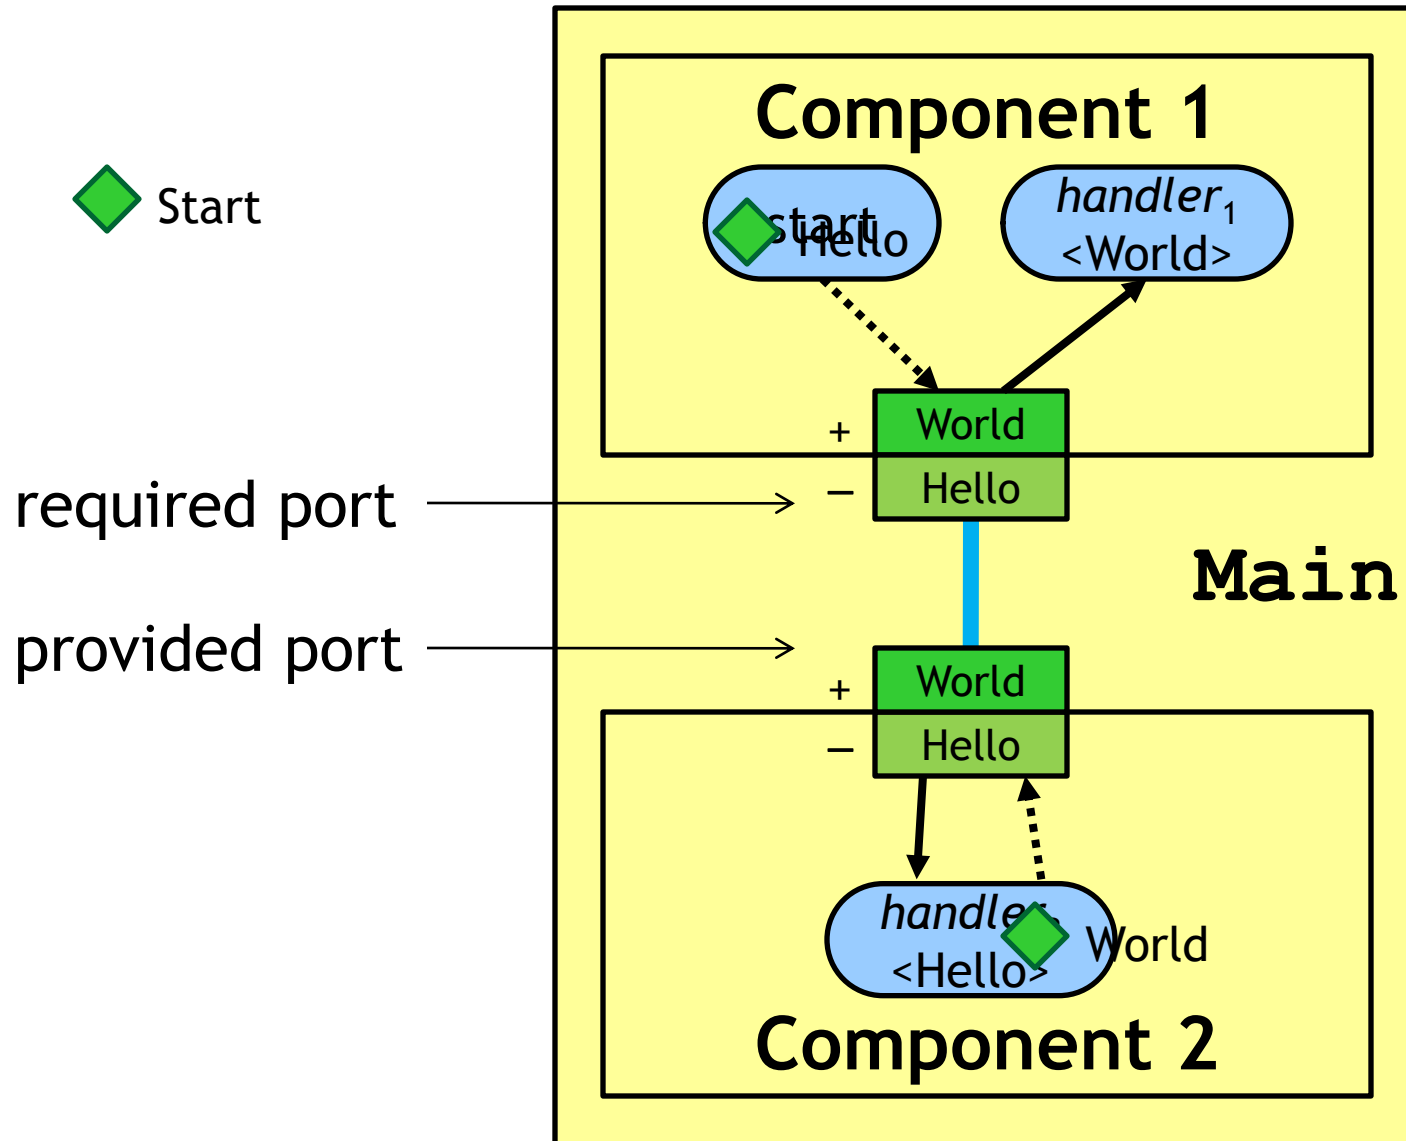


Components

Component

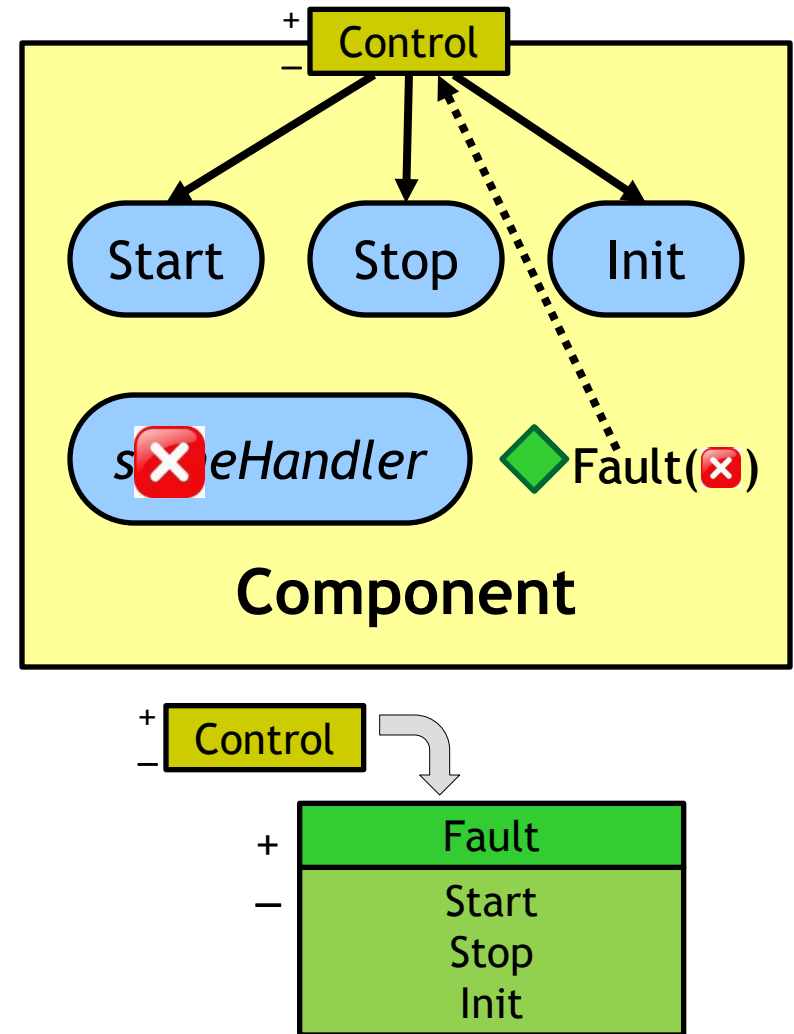
- are instantiated from component definitions
 - component definitions are Java classes
- A component instance is an object containing
 - local state variables
 - ports (provided or required interfaces)
 - event handlers
 - subscriptions
 - encapsulated subcomponents
 - channels
- form a containment hierarchy rooted at **Main**

Hello World! Example





Control port

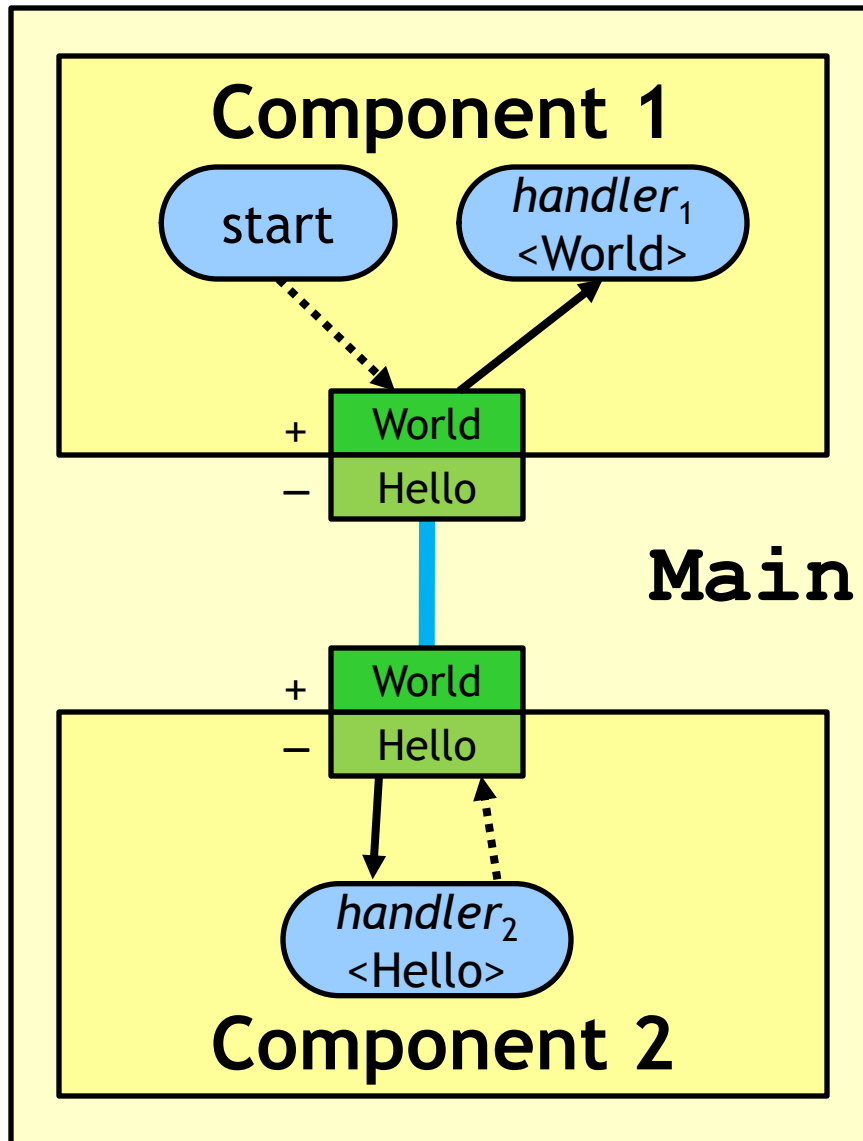
- Every component has a **Control** port
 - by default (provided)
 - not shown in diagrams
 - allows component to handle lifecycle events
- Exceptions / faults not caught inside a handler
 - wrapped into Fault event
 - triggered on control port



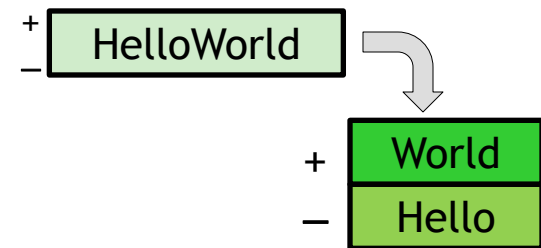
Init event

- Configures component with initial parameters
 - Specialized for each component definition
 - E.g. **MyComponent** handles **MyInit**  $\text{MyInit} \subseteq$  **Init**
- Init guaranteed to be the first event handled
 - “If a component subscribed an Init handler to its control port in its constructor, the component will not handle any other event before an Init event!”
 - The component will not do anything if the parent component does not trigger an Init event on its control port!

Hello World! Example



◆ Hello
◆ World



Source code: Events and Port

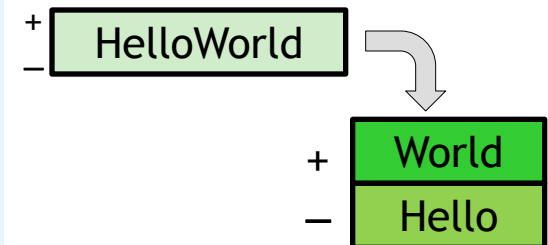
```
public final class Hello extends Event {  
    private final String message;  
    public Hello(String m) {  
        message = m;  
    }  
    public String getMessage() {  
        return message;  
    }  
}
```

◆ Hello

```
public final class World extends Event {  
}
```

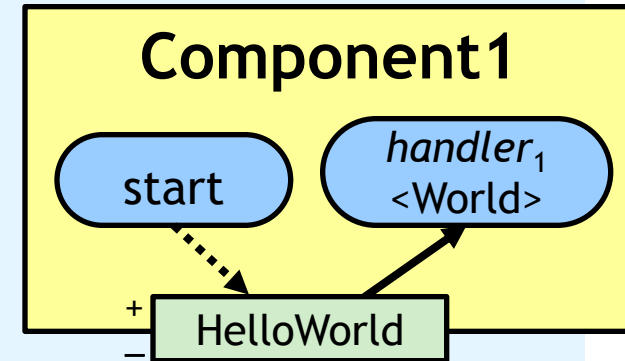
◆ World

```
public class HelloWorld extends PortType{  
    positive(World.class);  
    negative(Hello.class);  
}}
```



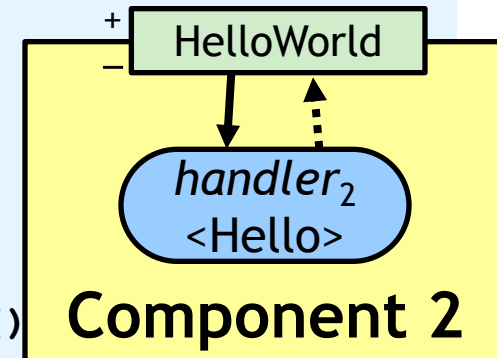
Source code: Component1

```
public class Component1 extends ComponentDefinition {  
  
    private Positive<HelloWorld> hwPort = positive(HelloWorld.class);  
  
    public Component1() {  
        System.out.println("Component1 created.");  
        subscribe(startHandler, control);  
        subscribe(worldHandler, hwPort);  
    }  
  
    Handler<Start> startHandler = new Handler<Start>() {  
        public void handle(Start event) {  
            System.out.println("Component1 started. Triggering Hello...");  
            trigger(new Hello("Hi there!"), hwPort);  
        }  
    };  
  
    Handler<World> worldHandler = new Handler<World>() {  
        public void handle(World event) {  
            System.out.println("Component1 received World event.");  
        }  
    };  
}
```



Source code: Component2

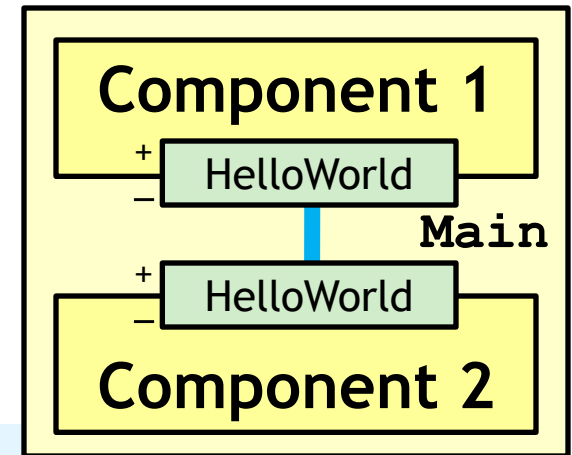
```
public class Component2 extends ComponentDefinition {  
    private Negative<HelloWorld> hwPort = negative(HelloWorld.class);  
    public Component2() {  
        System.out.println("Component2 created.");  
        subscribe(startHandler, control);  
        subscribe(helloHandler, hwPort);  
    }  
    Handler<Start> startHandler = new Handler<Start>() {  
        public void handle(Start event) {  
            System.out.println("Component2 started.");  
        }  
    };  
    Handler<Hello> helloHandler = new Handler<Hello>() {  
        public void handle(Hello event) {  
            System.out.println("Component2 received Hello event with "  
                               + "message: " + event.getMessage());  
            trigger(new World(), hwPort);  
        }  
    };  
}
```



Source code: Main

- Main is a Java main class

```
public class Main extends ComponentDefinition {
    private Component component1, component2;
    public Main() {
        System.out.println("Main created.");
        component1 = create(Component1.class);
        component2 = create(Component2.class);
        connect(component1.getNegative(HelloWorld.class),
                component2.getPositive(HelloWorld.class));
    }
    public static void main(String[] args) {
        Kompics.createAndStart(Main.class);
        Kompics.shutdown();
    }
}
```

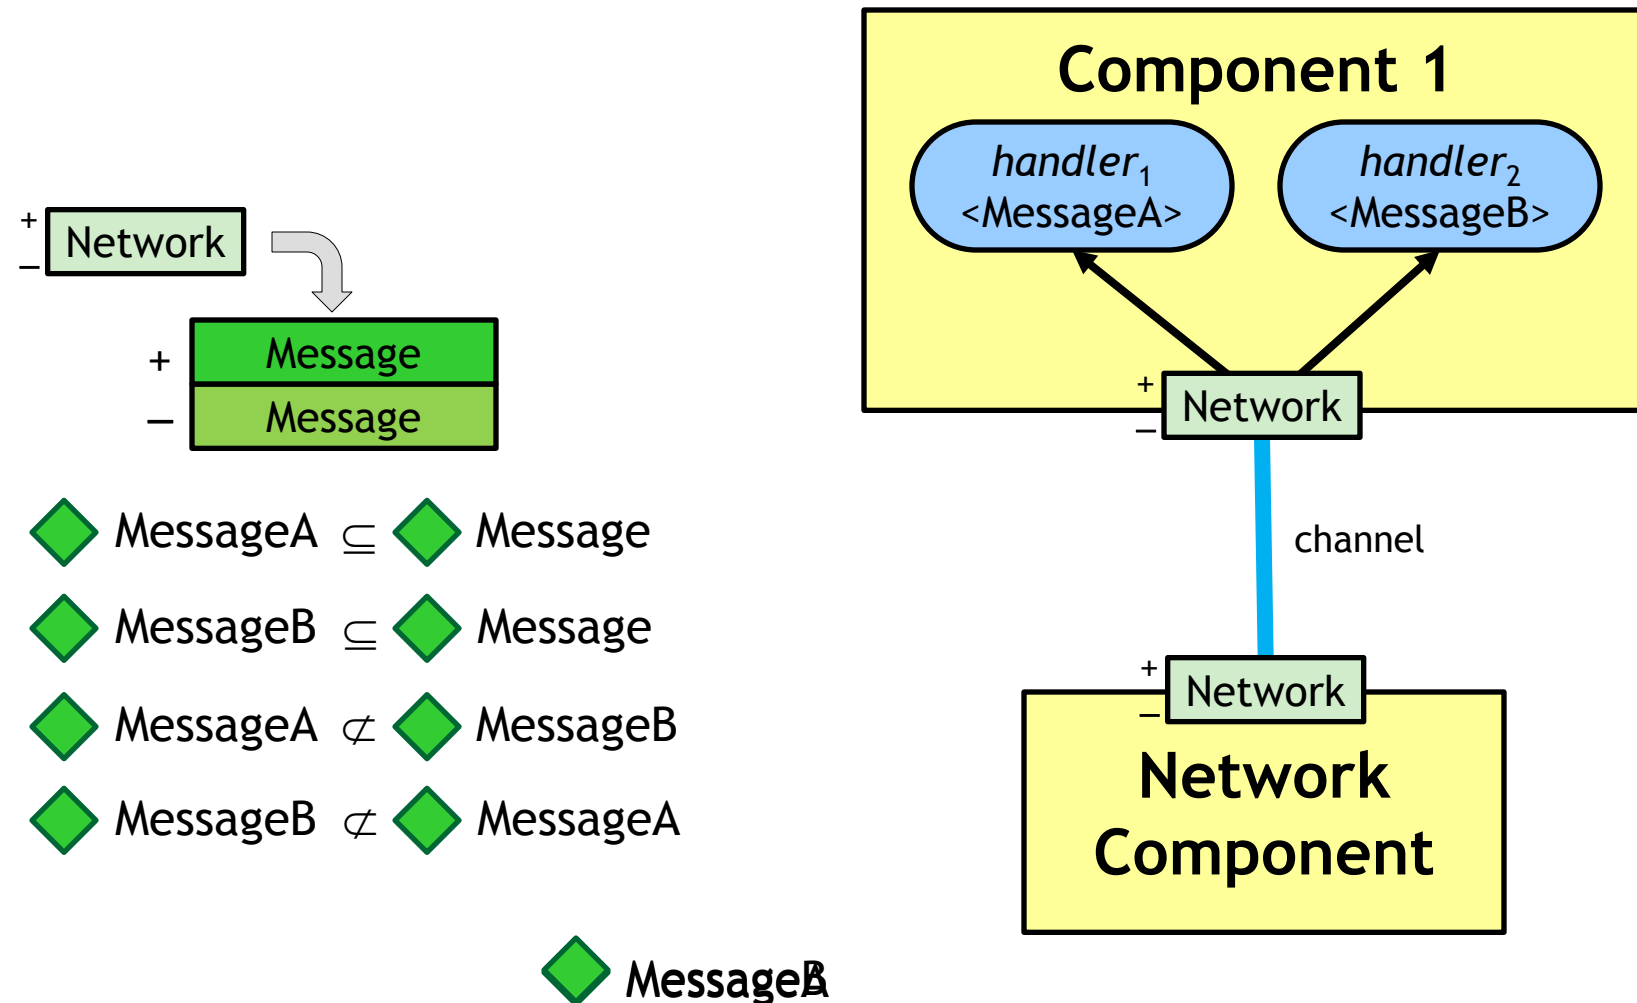


Output

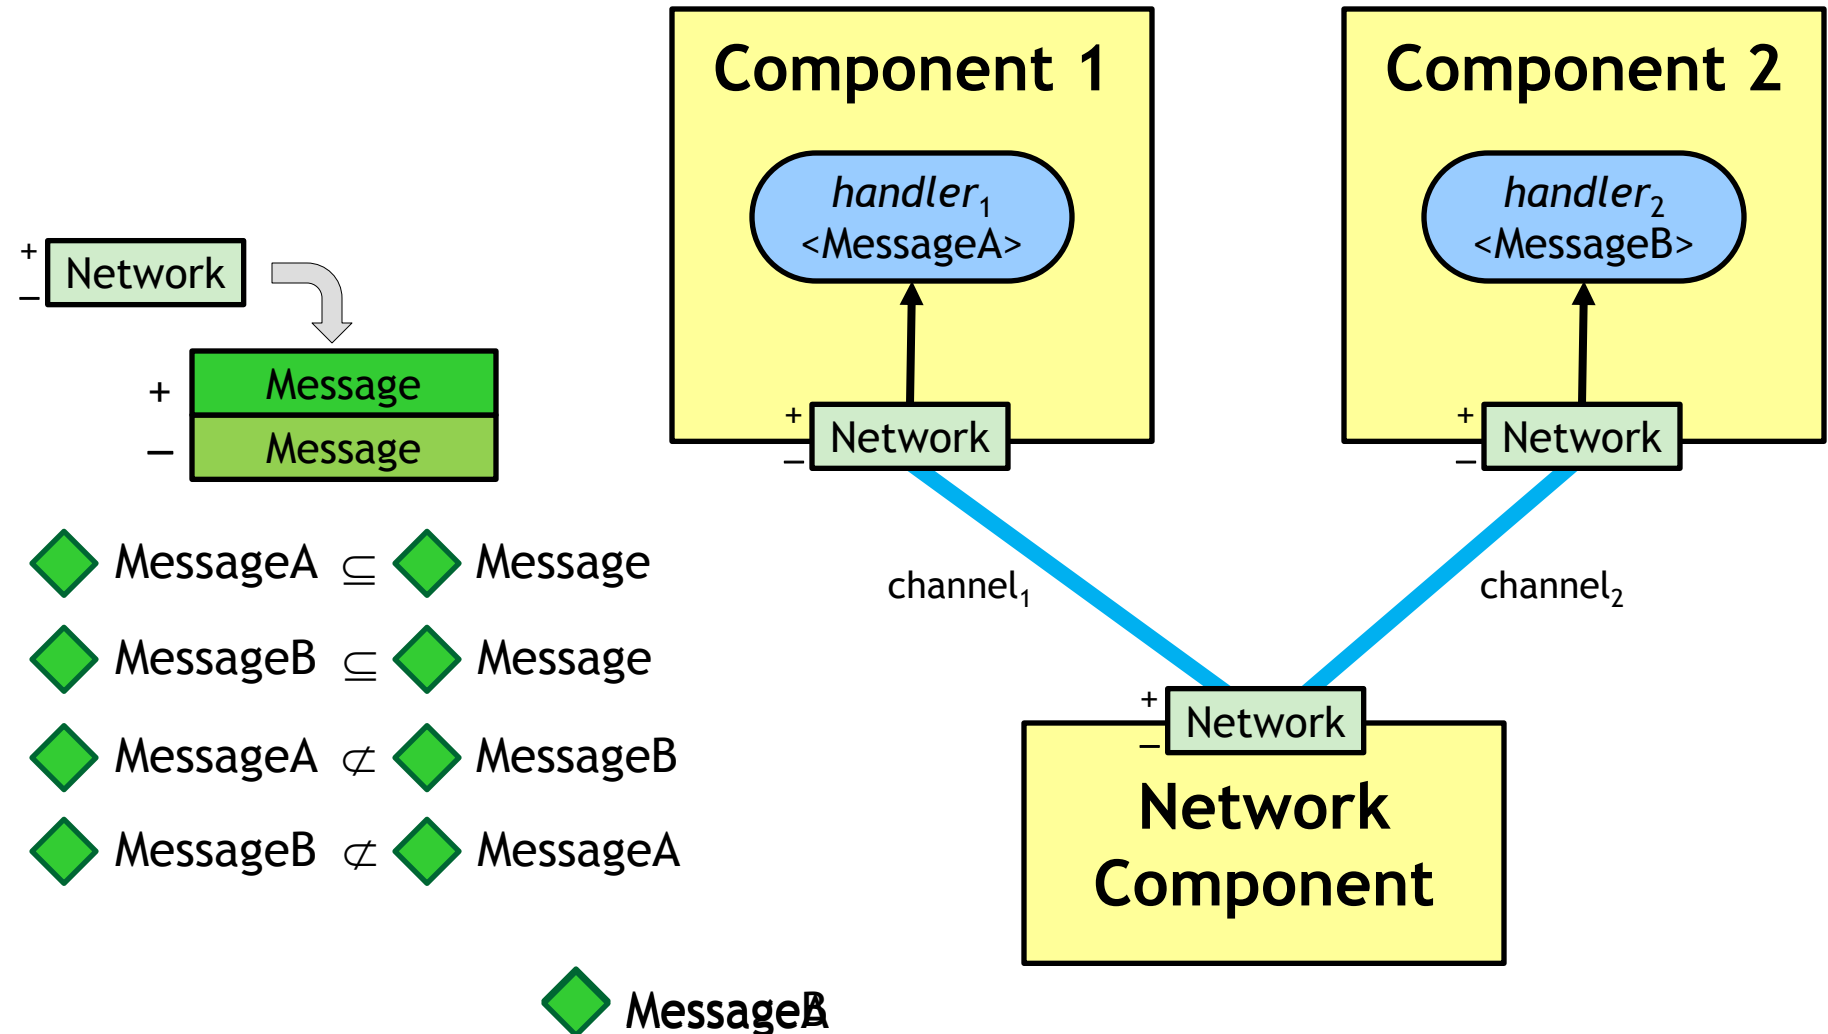
```
prompt:$ java Main
Main created.
Component1 created.
Component2 created.
Component1 started. Triggering Hello...
Component2 started.
Component2 received Hello event with message: Hi there!
Component1 received World event.
prompt:$ _
```

- Kompics runtime creates and starts **Main**
 - **Main** recursively creates and starts **c1**, **c2**
 - **c1**'s start handler triggers Hello event
 - handled by **c2**, which triggers World event
 - handled by **c1**

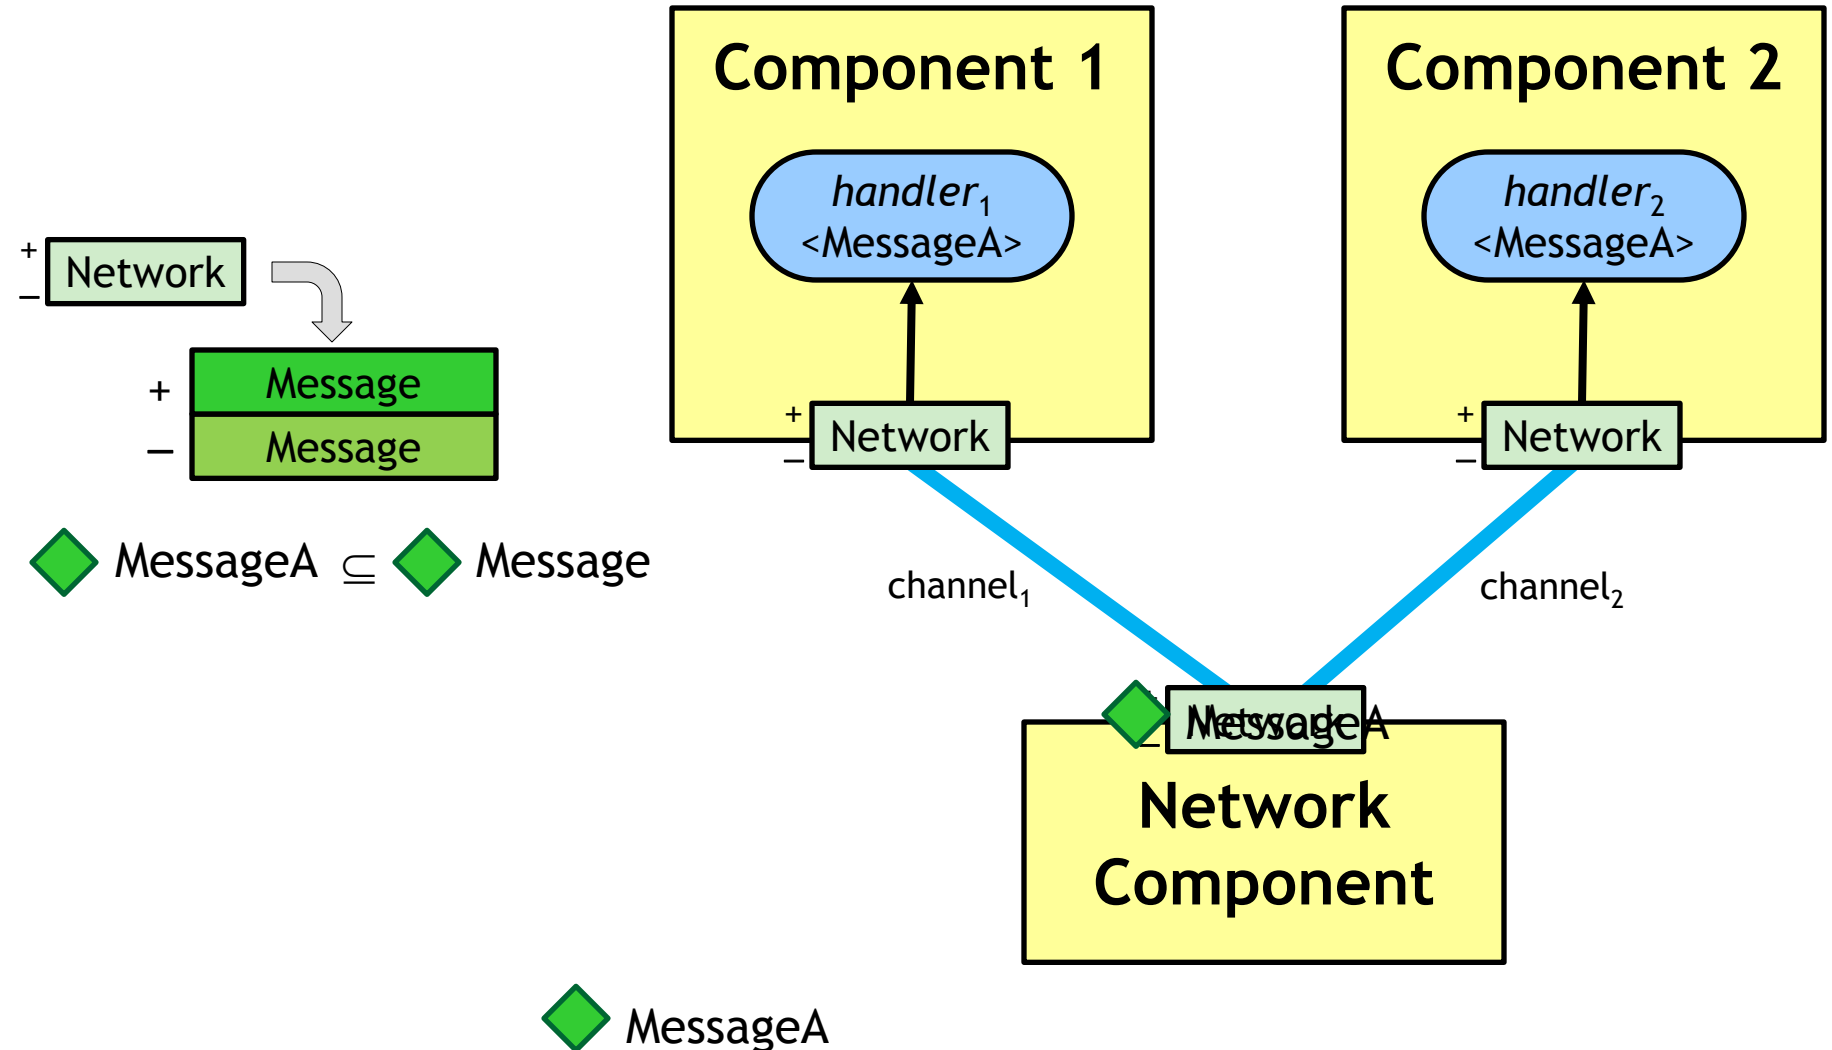
Publish / Subscribe



Publish / Subscribe

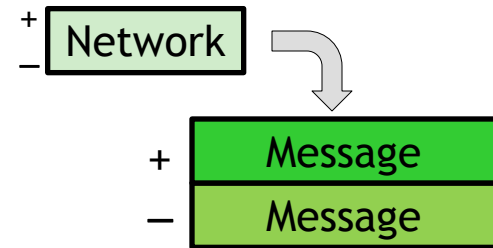


Publish / Subscribe



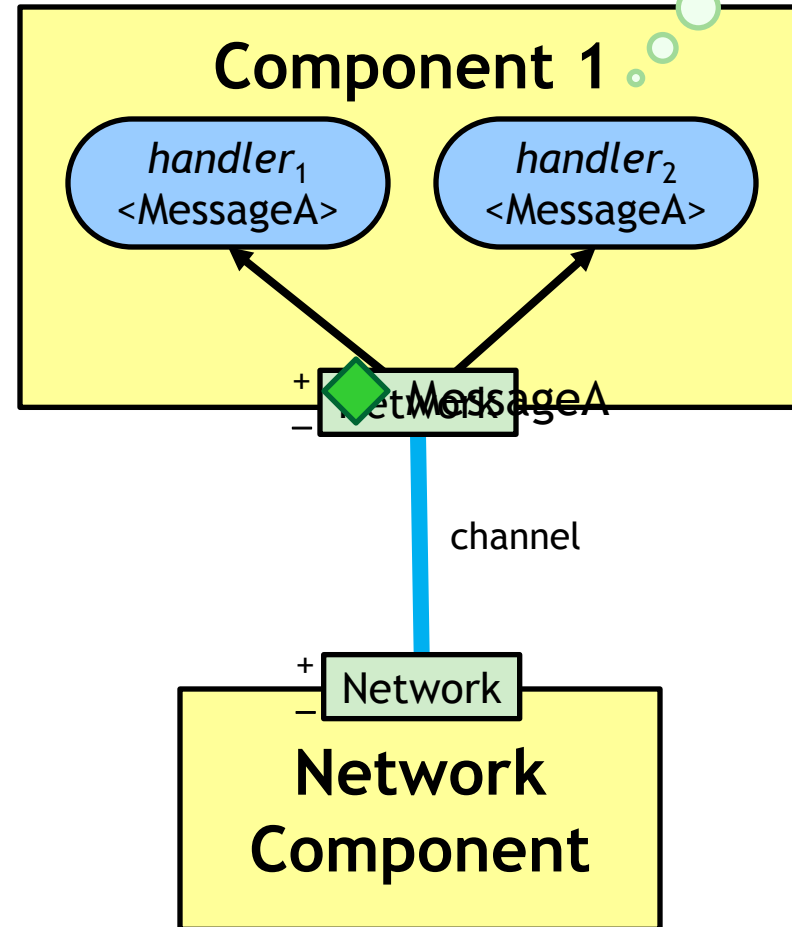
Publish / Subscribe

Sequential
handler
execution



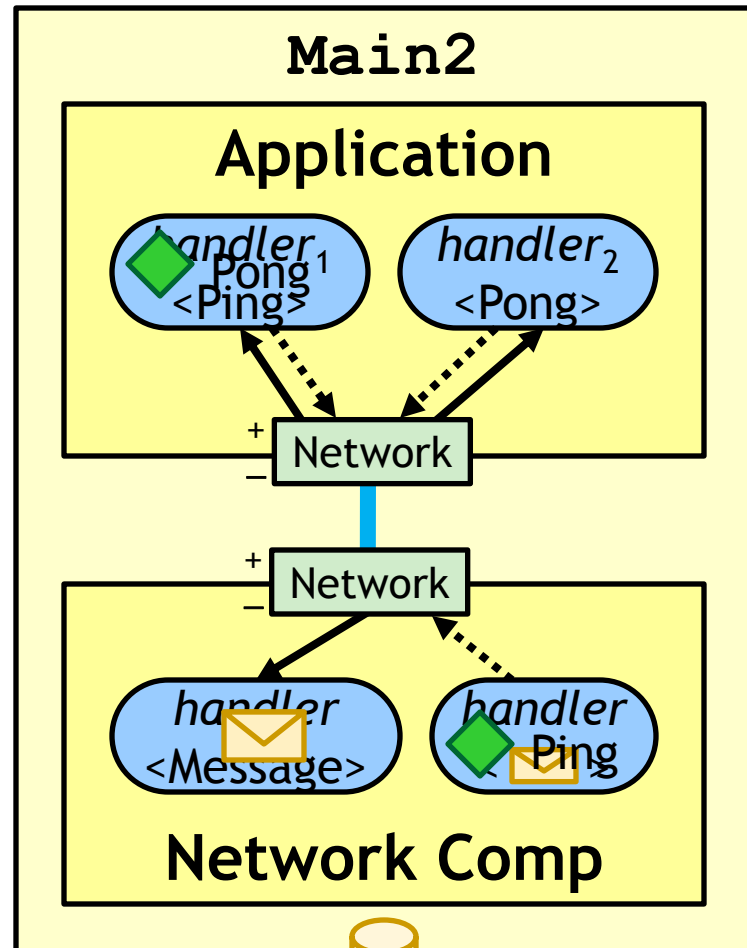
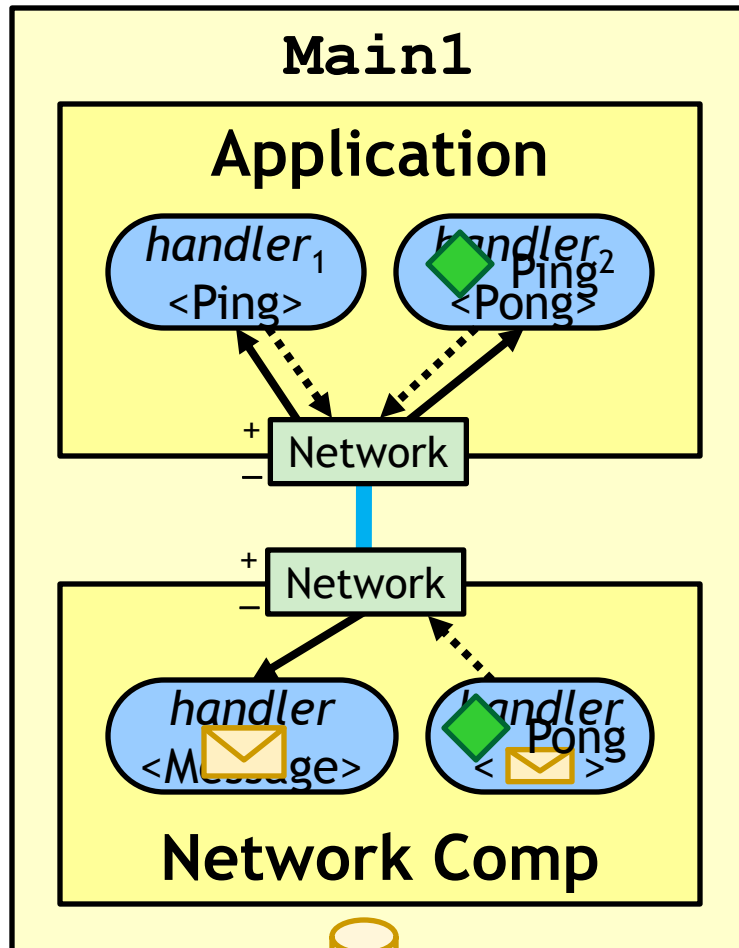
◆ MessageA \subseteq ◆ Message

◆ MessageA

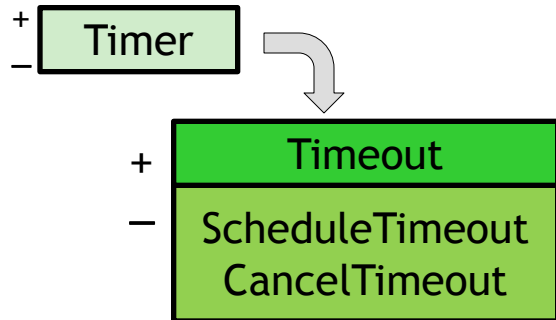


Send remote messages

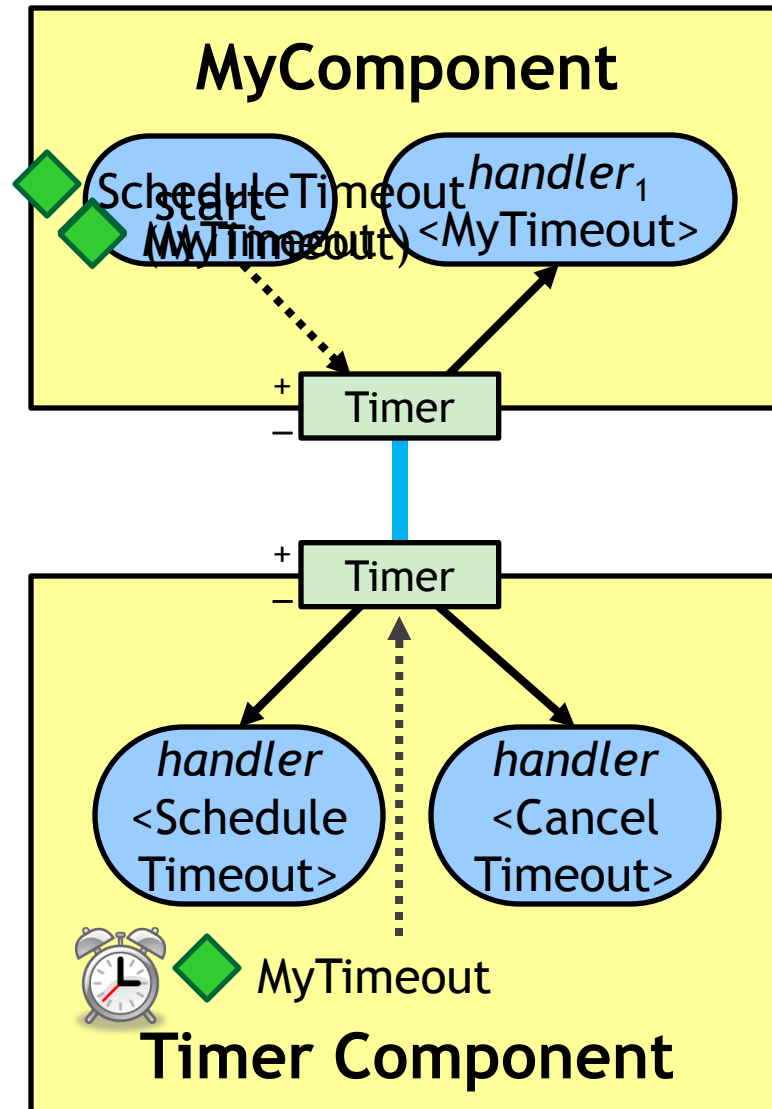
◆ Ping \subseteq ◆ Message
◆ Pong \subseteq ◆ Message



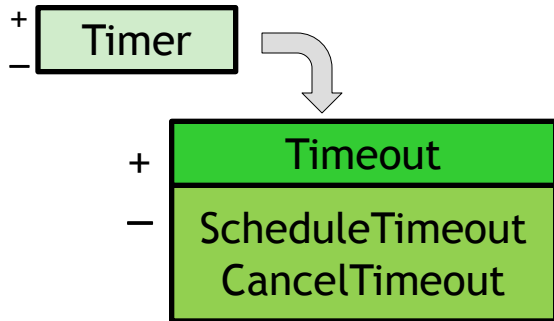
Scheduling a Timeout



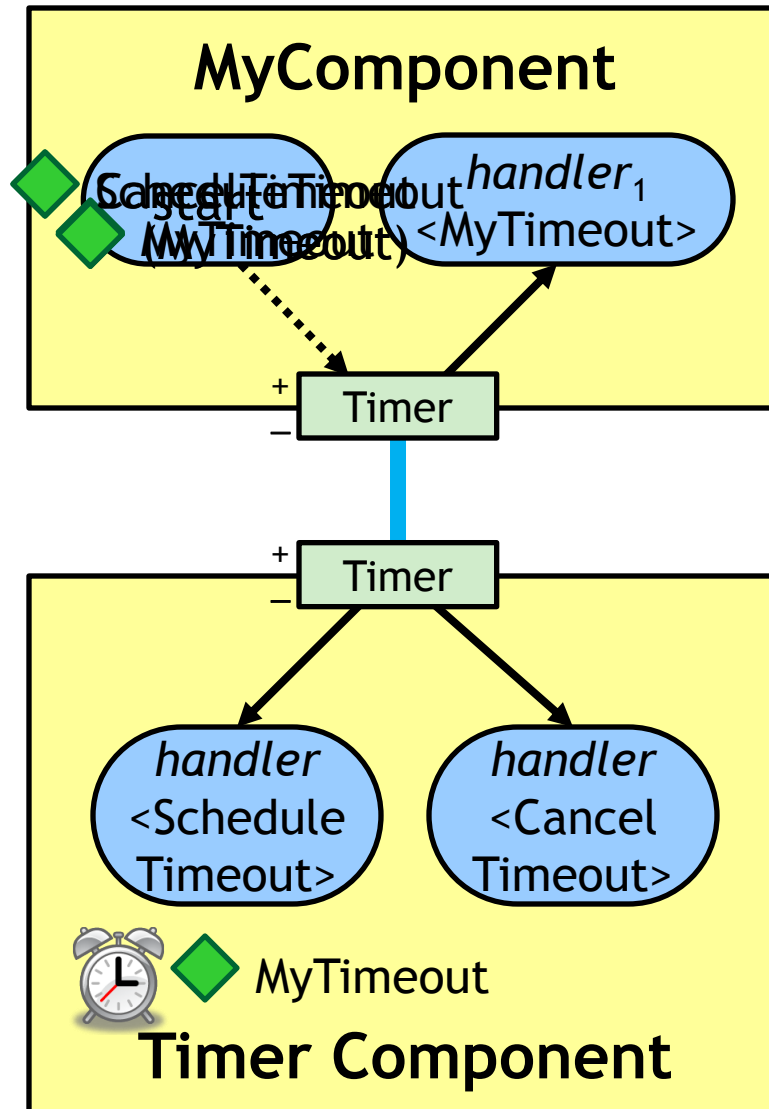
◆ `MyTimeout` \subseteq ◆ `Timeout`



Canceling a Timeout



◆ MyTimeout \subseteq ◆ Timeout



◆ MyTimeout

Timer Component

Scheduling & Canceling a Timeout

◆ MyTimeout \subseteq ◆ Timeout

```
class MyComponent extends ComponentDefinition {
    Positive<Timer> timer = positive(Timer.class); //required

    Handler<Start> startHandler = new Handler<Start>() {
        public void handle(Start event) {
            // scheduling a timeout
            long delay = 5000; // milliseconds
            ScheduleTimeout st = new ScheduleTimeout(delay);
            st.setTimeoutEvent(new MyTimeout(st));
            UUID timeoutId = st.getTimeoutEvent().getTimeoutId();
            trigger(st, timer);
            // canceling a timeout
            CancelTimeout ct = new CancelTimeout(timeoutId);
            trigger(ct, timer);
        }
    };
}
```

Software engineering view

- Events and ports are **interfaces**
 - service abstractions, modules
 - packaged together as libraries
- Components are **implementations**
 - provide or require modules / interfaces
 - dependencies on provided / required modules
 - expressed as library dependencies
 - multiple implementations for some module
 - separate libraries
 - deploy-time composition

Kompics software

- Open source project
- Latest release: 0.4.2.6
- <http://kompics.sics.se>
 - Java source code (SVN)
 - Javadocs
 - Test reports
- Ports and components (e.g. Network, Timer)
 - are packaged as JARs
 - You can use Maven to download them automatically
 - see <http://kompics.sics.se/trac/browser/trunk/pom.xml>

Relation to the textbook

How we use Kompics to model the
abstractions in the textbook / course

Correspondences

Textbook / Lectures Kompics

An **implemented** port is **positive** on the

A **required / used** port is **negative** on

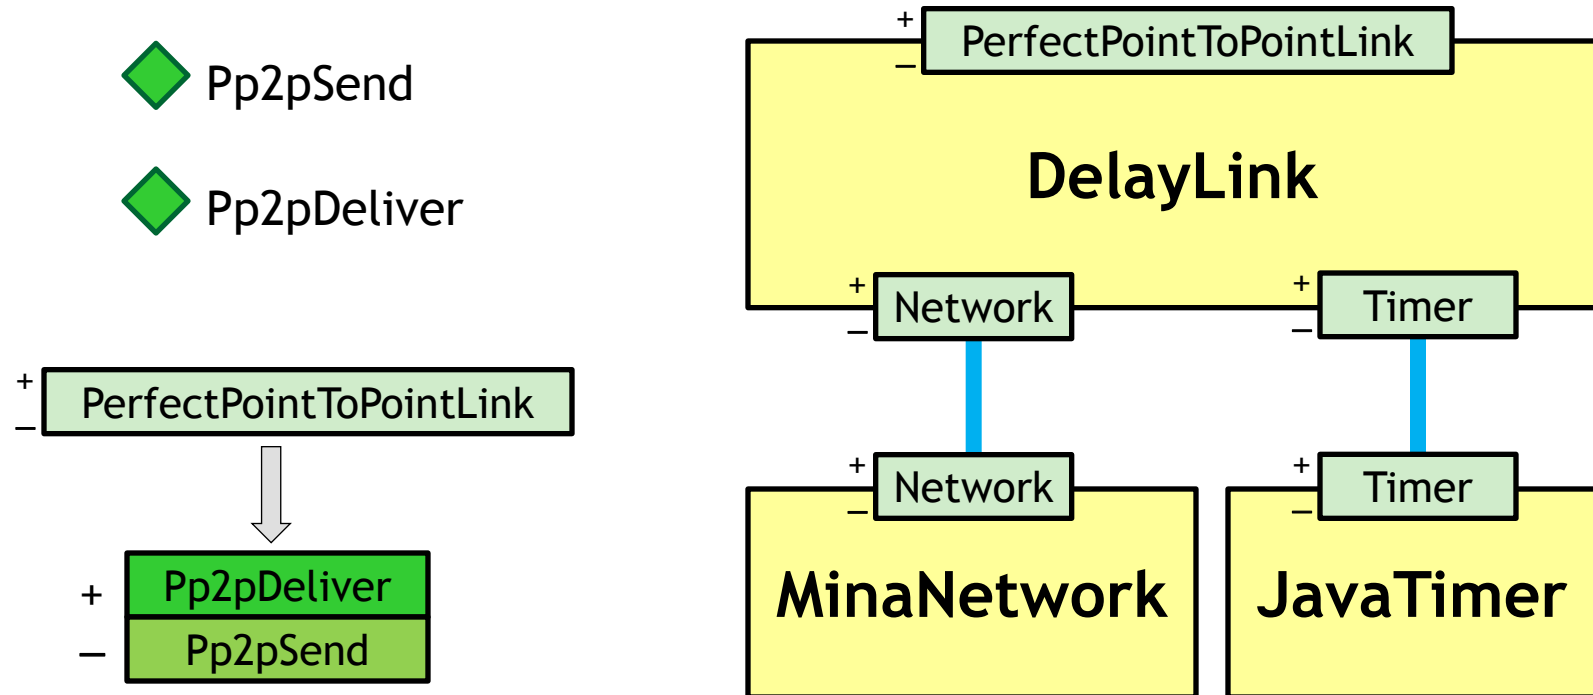
Ports

```
ScheduleTimeout st = new ScheduleTimeout( $\delta$ );  
st.setTimeoutEvent(new MyTimeout(st));  
trigger(st, timer);
```

on (**negative**, -)
(**positive**, +)

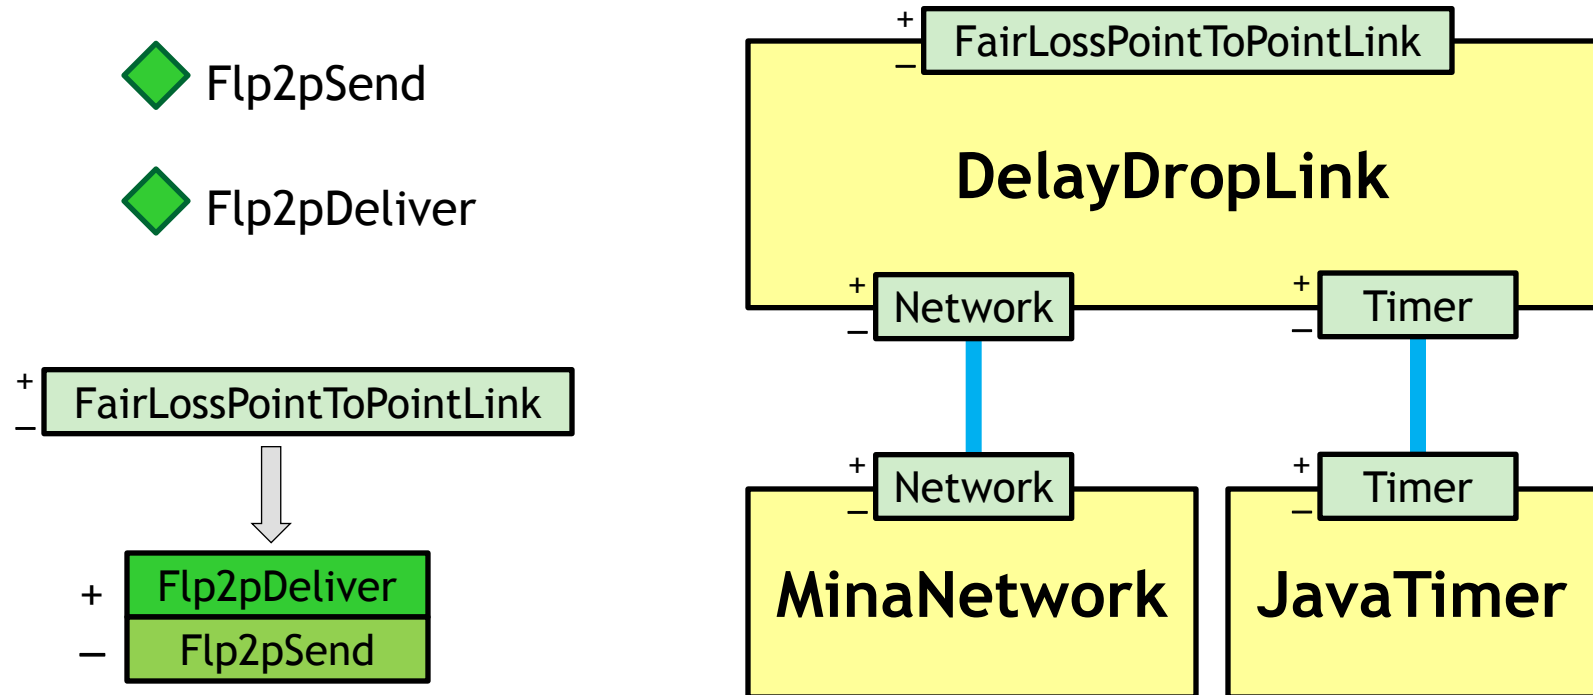
- Algorithms
 - Implements SomeModule
 - Uses SomeModule
 - *startTimer*(δ , MyTimeout);
- Components
- **Negative**<SomePort>
 - **Positive**<SomePort>
 - Trigger ScheduleTimeout event on Timer port

PerfectPointToPointLink



- Messages sent over perfect links should extend the Pp2pDeliver event

FairLossPointToPointLink



- Messages sent over fair-loss links should extend the Flp2pDeliver event

Assignments framework

How we run distributed systems and
create experiment scenarios

Experimentation framework

- Define network topologies
 - processes and their addresses: <id, IP, port>
 - properties of links between processes
 - latency (ms)
 - loss rate (%)
- Define execution scenarios
 - the sequence of service requests initiated by each process in the distributed system
- Experiment with various topologies/scenarios
 - Launch all processes locally on one machine

Distributed System Launcher

- Read complete documentation at
 - <http://kompics.sics.se/trac/wiki/DistributedSystemLauncher>

```
public final class Experiment1 {  
    public static final void main(String[] args) {  
        Topology topology1 = new Topology() {{  
            node(1, "127.0.0.1", 22031);  
            node(2, "127.0.0.1", 22032);  
            link(1, 2, 1000, 0).bidirectional();  
        }};  
        Scenario scenario1 = new Scenario(Main.class) {{  
            command(1, "S500:Lmsg1:S5000");  
            command(2, "S1000:Pmsg2:S4000");  
        }};  
        scenario1.executeOn(topology1);  
    }  
}
```

Process window

The image shows two windows from a system simulation. The top window, titled 'Process 1 - S500:Lmsg1:S5000', displays a terminal log with the following entries:

```
440@SCENARIO {Main} Process 1 has  
0 INFO {Application0} Sleeping 5  
615 INFO {Application0} Sending  
616 INFO {Application0} Sleeping  
2101 INFO {Application0} Recd  
5616 INFO {Application0} DONE
```

Below the log is an 'Input:' field. A callout bubble points to this field and says: 'After the Application completes the script it can process further commands input here...'. Another callout bubble points to the log and says: 'The script of service requests of the process is'.

The bottom window, titled 'Process 2 - S1000', displays a terminal log with the following entries:

```
446@SCENARIO {Main} Process 2 has [S1000:Pmsg2:S4000].  
0 INFO {Application0} Sleeping 4000 milliseconds...  
1144 INFO {Application0} Send select message msg2 to 1@127.0.0.1:22031  
1144 INFO {Application0} Sleeping 4000 milliseconds...  
1797 INFO {Application0} Received lossy message msg1  
5145 INFO {Application0} DONE ALL OPERATIONS
```

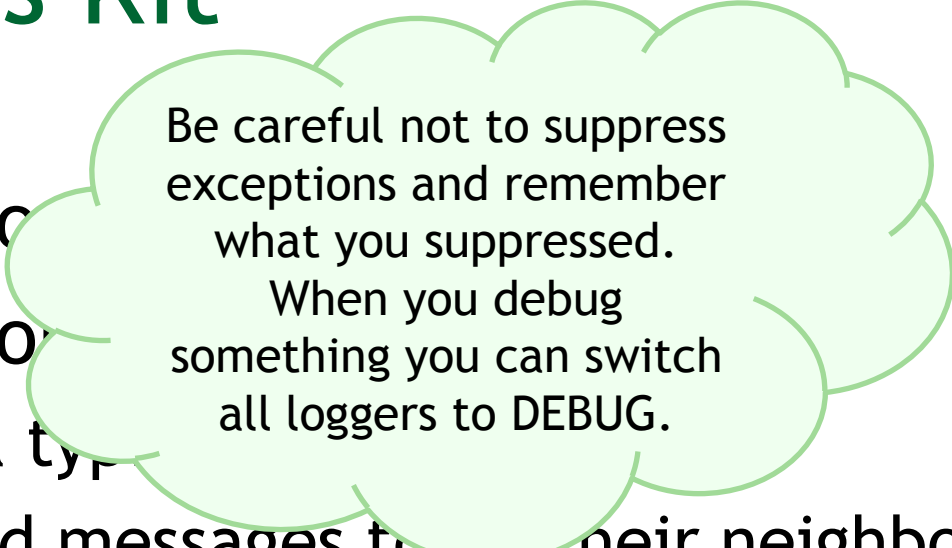
Below the log is an 'Input:' field. A callout bubble points to this field and says: 'Commands input in the right-hand input field are sent to all processes simultaneously...'. Another callout bubble points to the log and says: 'The script of service requests of the process is'.

To terminate an experiment press

Ctrl+K

Assignments Kit

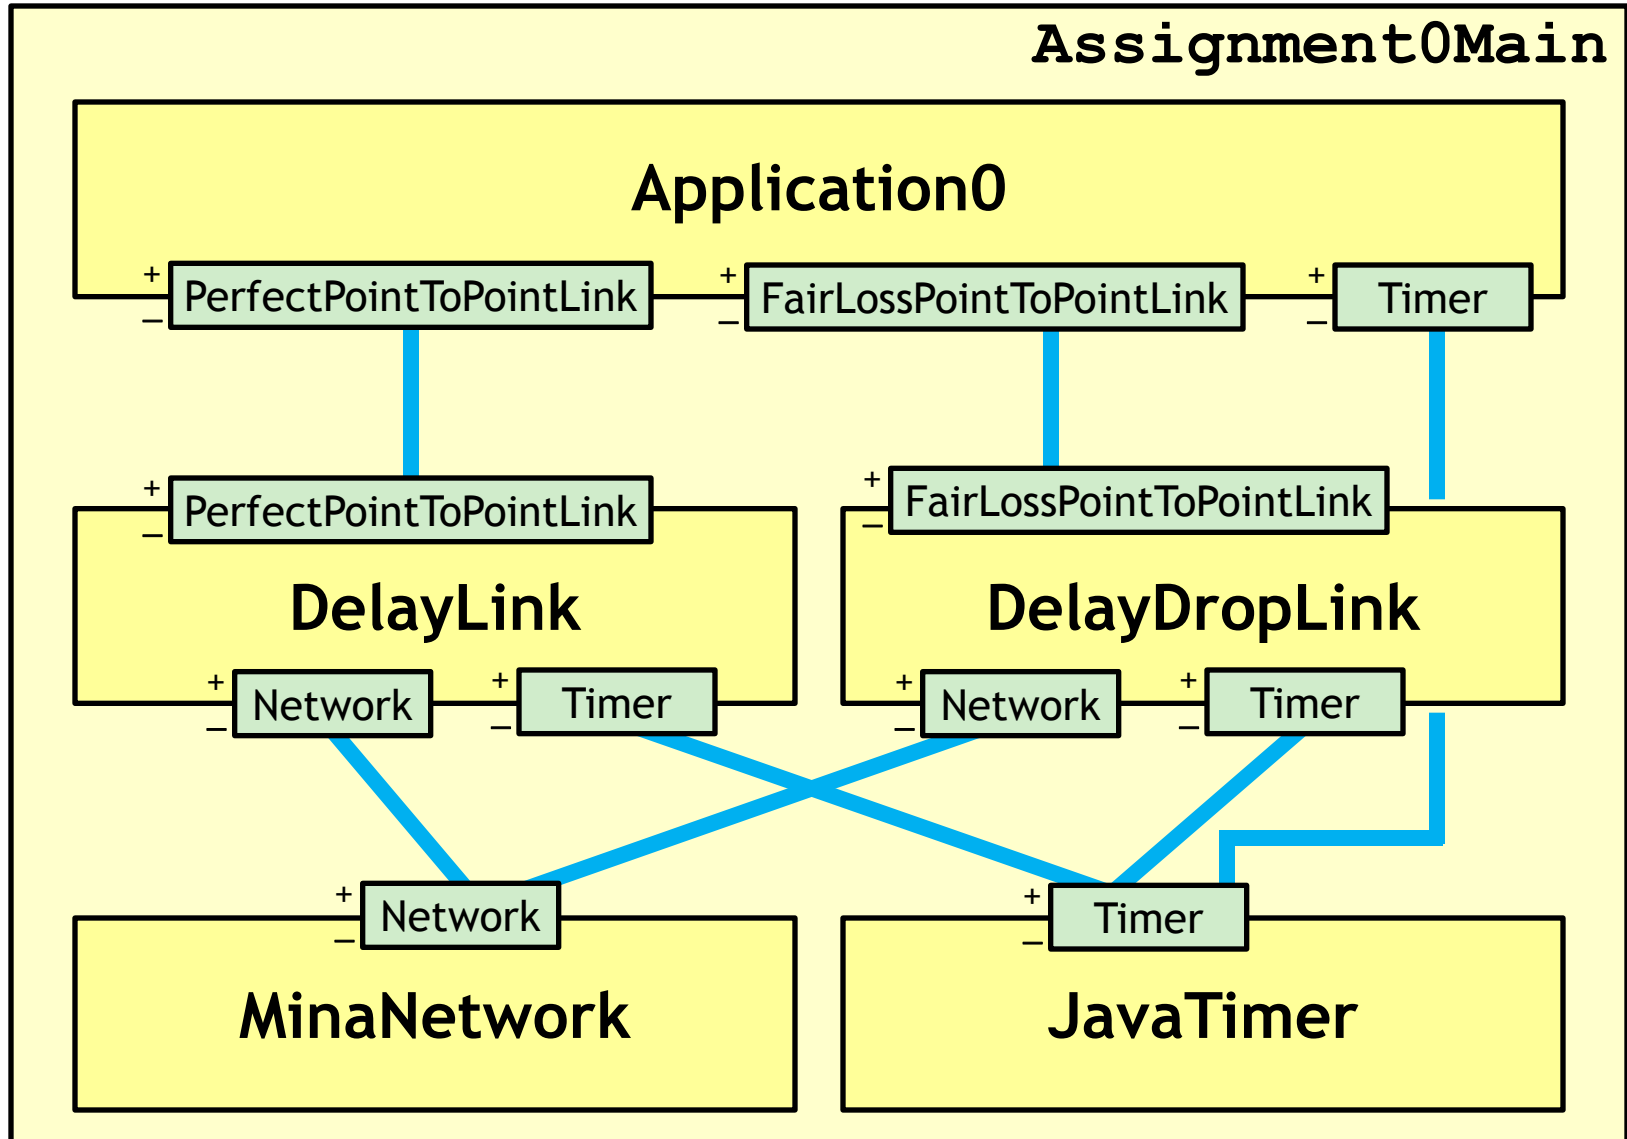
- Source code for
- Source code for
 - skeleton for a type
 - processes send messages to their neighbors
- Kompics and third party JARs
 - add them to your projects as libraries
 - `log4j.properties`: control logging output
- Miscellaneous
 - sample Eclipse project files
 - sample `pom.xml`, if you want to use Maven



Be careful not to suppress exceptions and remember what you suppressed.

When you debug something you can switch all loggers to DEBUG.

Architecture for Assignment0



First assignment

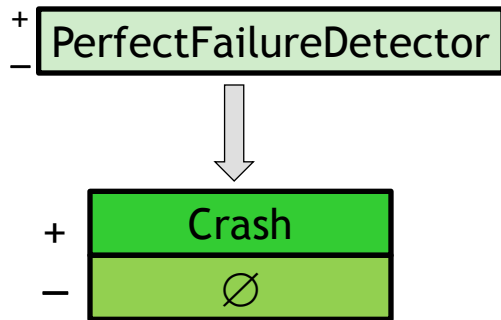
Failure detectors

First assignment: failure detectors

- Implement and experiment with
 - Perfect Failure Detector
 - Synchronous model
 - Eventually Perfect Failure Detector
 - Partially synchronous model
- Failure detectors discussed in lecture 4
- Read chapters 1 and 2 of the textbook!
- For requirements see Assignment1.pdf
- Use the algorithms in Assignment1.pdf!
 - Not the ones in the textbook!

Suggestions for PFD

◆ Crash

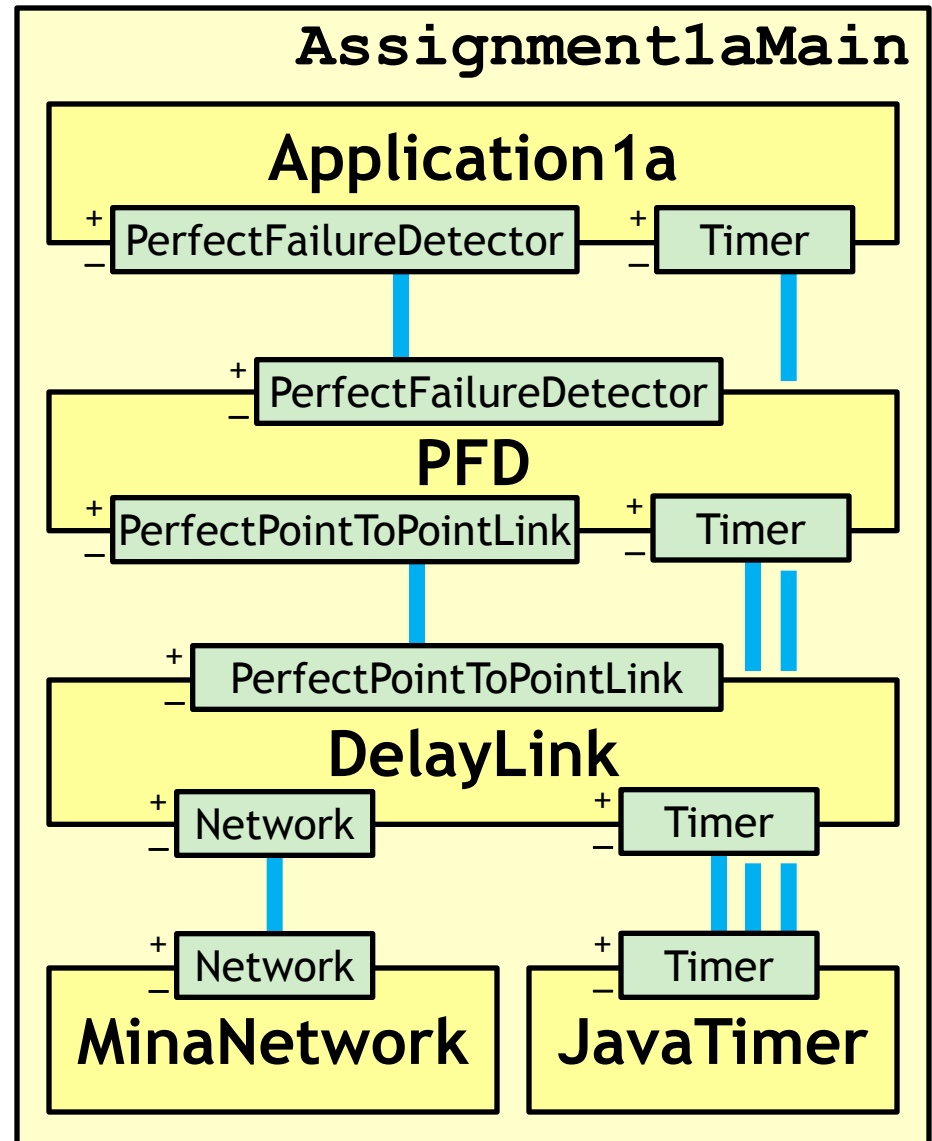


◆ CheckTimeout \subseteq ◆ Timeout

◆ HeartbeatTimeout \subseteq ◆ Timeout

◆ HeartbeatMessage \subseteq ◆ Pp2pDeliver

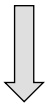
◆ PfdInit \subseteq ◆ Init



Suggestions for EPFD

◆ Suspect ◆ Restore

+ ◇PerfectFailureDetector
-

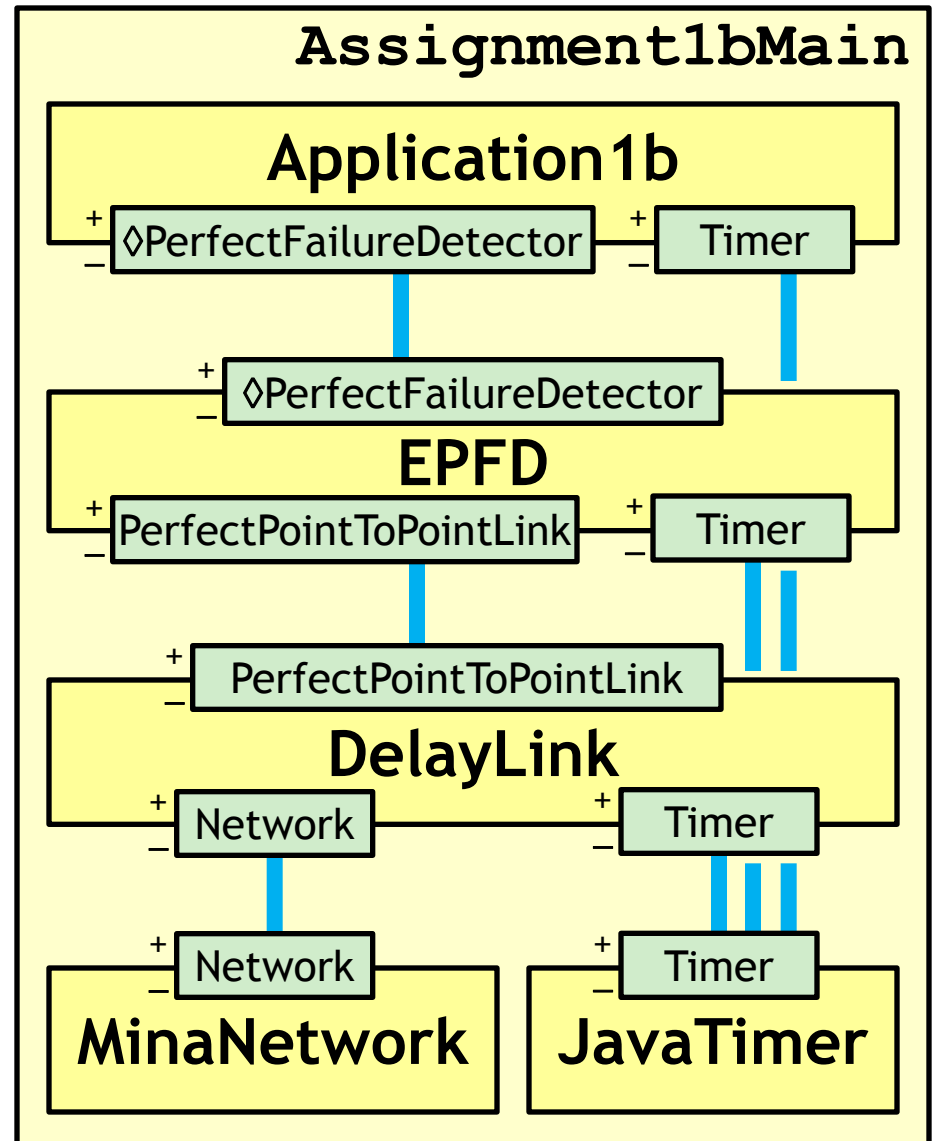


◆ CheckTimeout \subseteq ◆ Timeout

◆ HeartbeatTimeout \subseteq ◆ Timeout

◆ HeartbeatMessage \subseteq ◆ Pp2pDeliver

◆ EpfdInit \subseteq ◆ Init



Deliverables & Submission

- Deliverable 1: Source code ZIP archive
 - events, ports, components
 - exercise experiments (topologies and scenarios)
- Deliverable 2: PDF report
 - experiments observations and conclusions
 - answers to questions
- Submit deliverables to id2203.2010@gmail.com
 - Subject line: [assignmentX]LastName1LastName2
 - Before next tutorial starts (13:00 on deadline date)
 - Late submissions do not count

Presentations

- Present your implementation and experiments in class during next tutorial
- Both group members
 - Answer a few questions
 - You can present earlier (after some lecture)
 - You cannot present after the deadline

Forum

- Please use the forum to ask questions

<http://castor.sics.se/forum>

- Try to avoid direct emails



Good luck!

