
ID2203 - Tutorial 2

Distributed Systems, Advanced Course



Cosmin Arad

icarad(@)kth.se

Tallat M. Shafaat

tallat(@)kth.se

Seif Haridi

haridi(@)kth.se

KTH - The Royal Institute of Technology

This tutorial

- Broadcast primitives
 - Randomized/Probabilistic Broadcast
- Other flavours of broadcast taught in the class

Randomized algorithms

- Non-deterministic
- Probabilistic claims about guarantees
- Highly scalable and fault-tolerant
- Epidemics/gossiping

Interface of Probabilistic Broadcast

- **Module:**

- Name: ProbabilisticBroadcast (pb)

- **Events:**

- **Request:** $\langle \text{pbBroadcast} \mid m \rangle$
 - Broadcast m to all processes
- **Indication:** $\langle \text{pbDeliver} \mid \text{src}, m \rangle$
 - Deliver message m from process src

- **Properties:**

- ***PB1: Probabilistic Validity***
- ***PB2: No duplication***
- ***PB3: No creation***

Implementing pb

- Lazy probabilistic broadcast
- Main idea:
 - Use a cheap **unreliable** broadcast to disseminate a message
 - Some nodes randomly **save** the message
 - Use gossiping to **recover** any lost messages
- How to know if a message is lost?
 - Sender tags messages with sequence numbers
 - Gap in sequence numbers of received messages indicates loss of messages

Implementing pb

■ Init

- $lsn := 0$ ► *local sequence number*
- $del[\forall p] := 0$ ► *sequence number of last message delivered sent by p_i*
- $pending := \{\}$ ► *set of undelivered messages*

■ $\langle pbBroadcast \mid m \rangle$

- $lsn := lsn + 1$
- $\langle unBroadcast \mid [m, lsn] \rangle$ ► *unreliably broadcast m with sequence number*

Implementing pb

- $\langle \text{unreliableDeliver} \mid p, m, sn \rangle$
 - Randomly decide if I should store m
 - If I have already delivered ' $sn-1$ ' from p
 - **trigger** $\langle \text{pbDeliver} \mid m \rangle$
 - Move delivered pointer to sn
 - If I have missed some message(s)
 - Add m to the *pending* set
 - For all missed messages, initiate gossip
 - Wait for a 'timeout' for the gossip to succeed

Implementing pb

- $\langle \text{Timeout} \mid p, sn \rangle$
 - Deliver as many messages as received
 - Move delivered pointer to sn

Suggestions for lpb

◆ SUBInit \subseteq ◆ Init

◆ LazyPB \subseteq ◆ Init

◆ pbBroadcast ◆ pbDeliver

◆ unBroadcast ◆ unDeliver

◆ LPBTimeout \subseteq ◆ Timeout

◆ RequestMessage \subseteq ◆ Flp2pDeliver

◆ DataMessage \subseteq ◆ Flp2pDeliver

+ ProbabilisticBroadcast
-



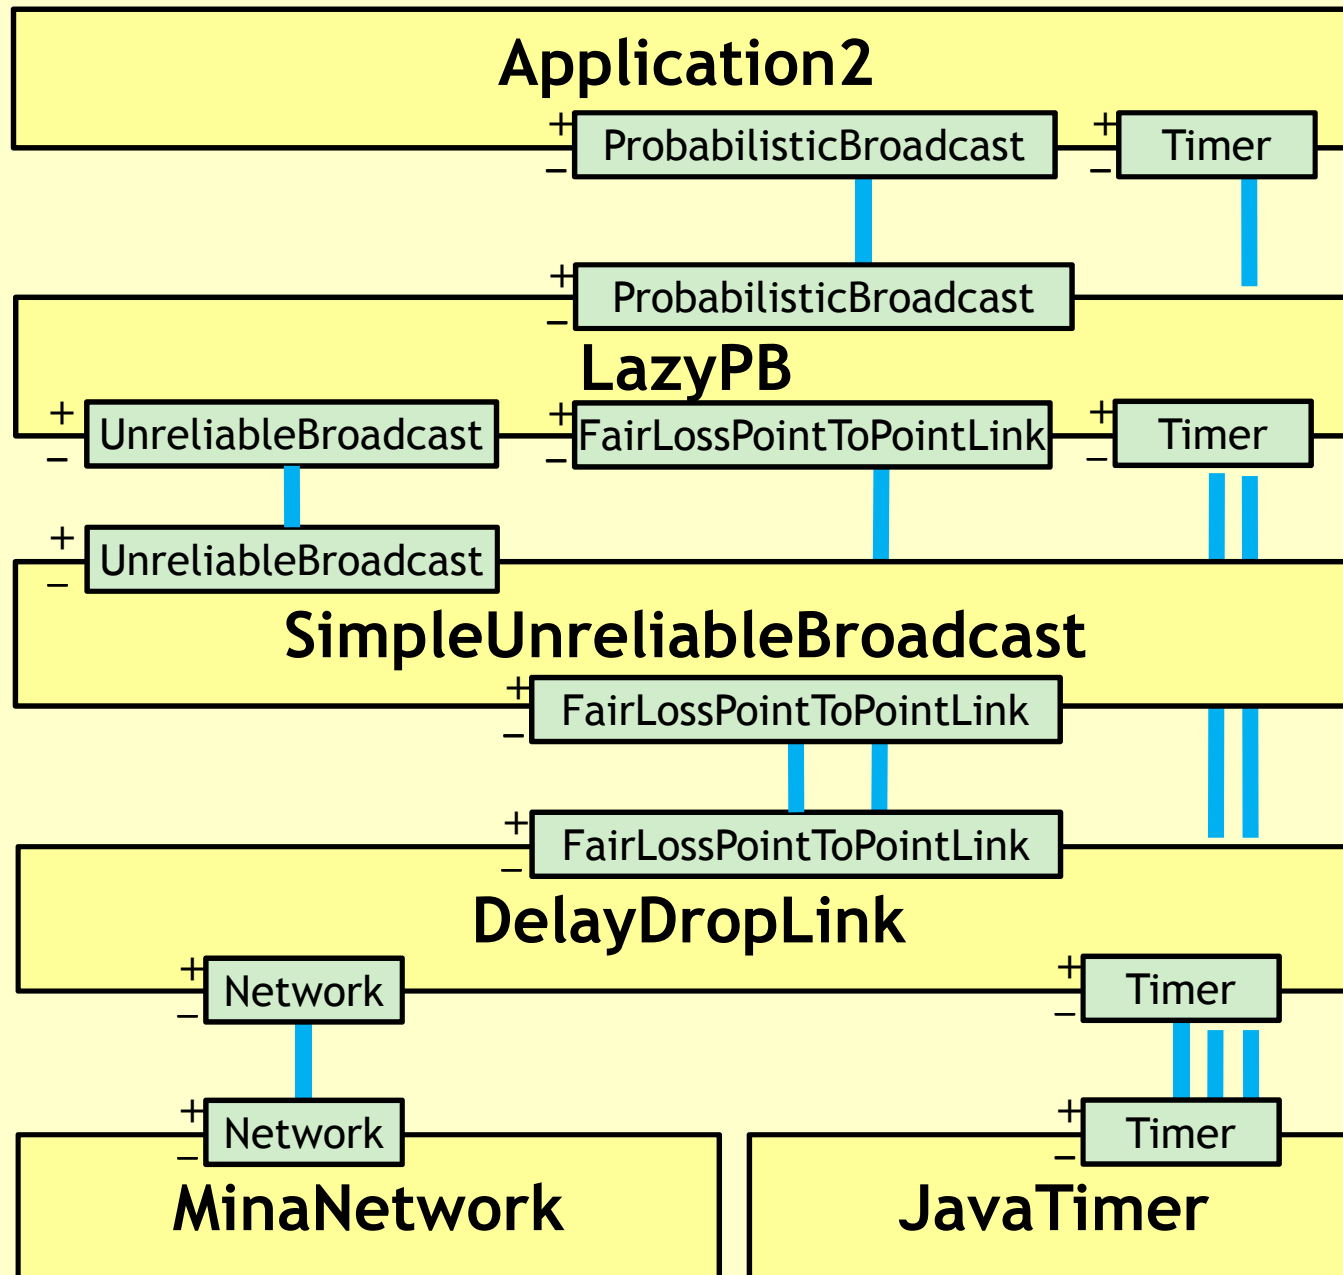
+ pbBroadcast
- pbDeliver

+ UnreliableBroadcast
-



+ unBroadcast
- unDeliver

Assignment2Main



Exercises

■ Exercise 1

- ❑ Correct the algorithm.
- ❑ Give your new algorithm in your report along with discussion on the errors fixed.
- ❑ Implement your algorithm in Kompics.
- ❑ The algorithm depends on *fanout*, *store-threshold* and *maxrounds*. Discuss their effect on the broadcast.

Exercises

■ Exercise 2

- Lazy probabilistic broadcast assumes you have a fully connected topology. Discuss what will happen if the topology is not fully connected. Discuss any modifications needed.
 - You do not need to implement this!

- As always, questions on the forum are welcome!