

Probabilistic Broadcast

Course leader: Professor Seif Haridi


Assistants: **Cosmin Arad**, Tallat Shafaat

{haridi, icarad, tallat}@kth.se


Motivation

- As systems grow larger and larger
 - Deterministic reliability is more and more costly
- Solution
 - Trade deterministic reliability guarantees for **scalability**
 - Provide **probabilistic reliability** guarantees
- Popular implementation technique
 - Epidemic dissemination (aka **gossip**)


Epidemic dissemination




Epidemic dissemination, round 0




Epidemic dissemination, round 1



Epidemic dissemination, round 2



Epidemic dissemination, round 3



Epidemic dissemination Pro&Con

■ Advantages

- ❑ Fast dissemination (low number of rounds)
- ❑ Cheap (constant # of messages per round)
- ❑ Nodes are equally loaded (no bottlenecks)
- ❑ Simple (compared with building a tree)

■ Disadvantaged

- ❑ Redundancy
- ❑ Not guaranteed to reach all nodes

Epidemic dissemination params

■ Fanout

- # of nodes contacted by one node in each round
- 2 in our previous example

■ Max rounds

- Maximum number of rounds for which a message is going to be retransmitted

Probabilistic broadcast (pb)

■ *Events*

- Request: $\langle \text{pbBroadcast} \mid m \rangle$
- Indication: $\langle \text{pbDeliver} \mid \text{src}, m \rangle$

■ *Properties: PB1, PB2, PB3*

- *PB1: Probabilistic Validity*
- *PB2: No duplication*
- *PB3: No creation*

Probabilistic broadcast (pb)

- ***Intuitively:*** messages are delivered with a certain probability
- ***Properties***
 - ***Probabilistic Validity:*** There is a given non-zero probability such that for any two correct processes p_i and p_j , every message broadcast by p_i is **eventually** delivered by p_j with this **probability**
 - ***No duplication:*** No message delivered more than once
 - ***No creation:*** No message delivered unless broadcast

Algorithms

- Eager probabilistic broadcast
- Lazy probabilistic broadcast

Eager probabilistic broadcast

- Main idea:
 - Similar to our previous example
 - Upon *pbBroadcast*, node picks *fanout* targets
 - Distinct and different from itself
 - Sends the message to targets
 - Message carries a time-to-live (TTL)
 - Initially $TTL = maxrounds$
 - May use fair-loss links
 - Receivers deliver filtering duplicates
 - They forward the message with a $TTL_{new} = TTL_{msg} - 1$
 - To *fanout* distinct and different random nodes...

Lazy probabilistic broadcast

- Main idea:
 - Use a cheap **unreliable** broadcast to disseminate a message
 - Some nodes **randomly save** the message
 - Use gossiping to **recover** any lost messages
- How to know if a message is lost?
 - Sender tags messages with **sequence** numbers
 - Gap in sequence numbers of received messages indicates loss of messages